

Decision Trees

Kaiqi Zhao
The University of Auckland

Slides are partially based on the materials from the University of British Columbia

Content

- Machine Learning Example
- Decision Trees
- Hypothesis Space
- Decision Tree Learning
- Supervised Learning
- Unsupervised Learning

Example: Should we wait for a table?

- You want to figure out whether you should wait for a table or not in a restaurant.

	Alt	Bar	Fri	Patrons	Price	Res	Type	Est	WillWait
x_1	Yes	No	No	Some	\$\$\$	Yes	French	Short	Yes
x_2	Yes	No	No	Full	\$	No	Thai	Mid	No
x_3	No	Yes	No	Some	\$	No	Burger	Short	Yes
x_4	Yes	No	Yes	Full	\$	No	Thai	Mid	Yes
x_5	Yes	No	Yes	Full	\$\$\$	Yes	French	Long	No
x_6	No	Yes	No	Some	\$\$	Yes	Italian	Short	Yes

[Example adapted from the AI book of Stuart J. Russell and Peter Norvig]

Alt: Is there any other suitable restaurants nearby?

Bar: Is there any bar area to wait in?

Fri: Is it Friday or Saturday?

Patrons: How many people are there (Some, Full)

Price: price range (\$, \$\$, \$\$\$)

Res: reservation of table?

Type: the type of restaurant (French, Thai, Burger, Italian)

Est: estimated waiting time (Short, Mid, Long)

Example: Should we wait for a table?

- You want to figure out whether you should wait for a table or not in a restaurant.

	Alt	Bar	Fri	Patrons	Price	Res	Type	Est	WillWait
x_1	Yes	No	No	Some	\$\$\$	Yes	French	Short	Yes
x_2	Yes	No	No	Full	\$	No	Thai	Mid	No
x_3	No	Yes	No	Some	\$	No	Burger	Short	Yes
x_4	Yes	No	Yes	Full	\$	No	Thai	Mid	Yes
x_5	Yes	No	Yes	Full	\$\$\$	Yes	French	Long	No
x_6	No	Yes	No	Some	\$\$	Yes	Italian	Short	Yes

[Example adapted from the AI book of Stuart J. Russell and Peter Norvig]

- How can we find the pattern and determine what leads to the decision?
 - Can you look at one attribute at a time?
 - Should we look at all attributes?

A Naïve Model

Example	Features (attributes)								Class labels
	Alt	Bar	Fri	Patrons	Price	Res	Type	Est	
x_1	Yes	No	No	Some	\$\$\$	Yes	French	Short	
x_2	Yes	No	No	Full	\$	No	Thai	Mid	
x_3	No	Yes	No	Some	\$	No	Burger	Short	
x_4	Yes	No	Yes	Full	\$	No	Thai	Mid	
x_5	Yes	No	Yes	Full	\$\$\$	Yes	French	Long	
x_6	No	Yes	No	Some	\$\$	Yes	Italian	Short	

⇒

WillWait
Yes
No
Yes
Yes
No
Yes

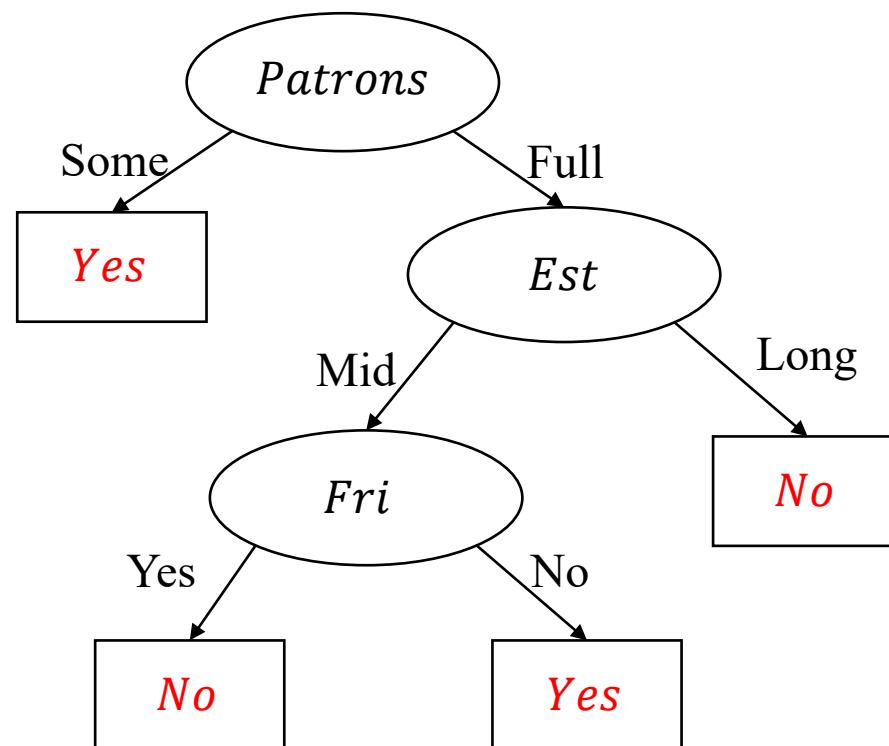
- A very naïve model – always predict one class label
 - Count how many times each label occurred in the data (4 Yes vs. 2 No)
 - Always predict the most common label, e.g., Yes.
- **Problems:** Not accurate! Features are not considered!
- How to leverage the features?

Decision Trees

Decision Trees

- Decision trees are simple models consisting of
 - A nested sequence of "if-else" decisions based on the features (splitting rules)
 - A **class label** as a return value at the end of each sequence

Can draw sequences of decisions as a tree



Example decision tree

```
if patrons = Some then
    return Yes
Else if patrons = Full then
    if Est = Long then
        return No
    else
        if Fri = Yes then
            return No
        else
            return Yes
    end
end
```

Another Example - Food Allergies

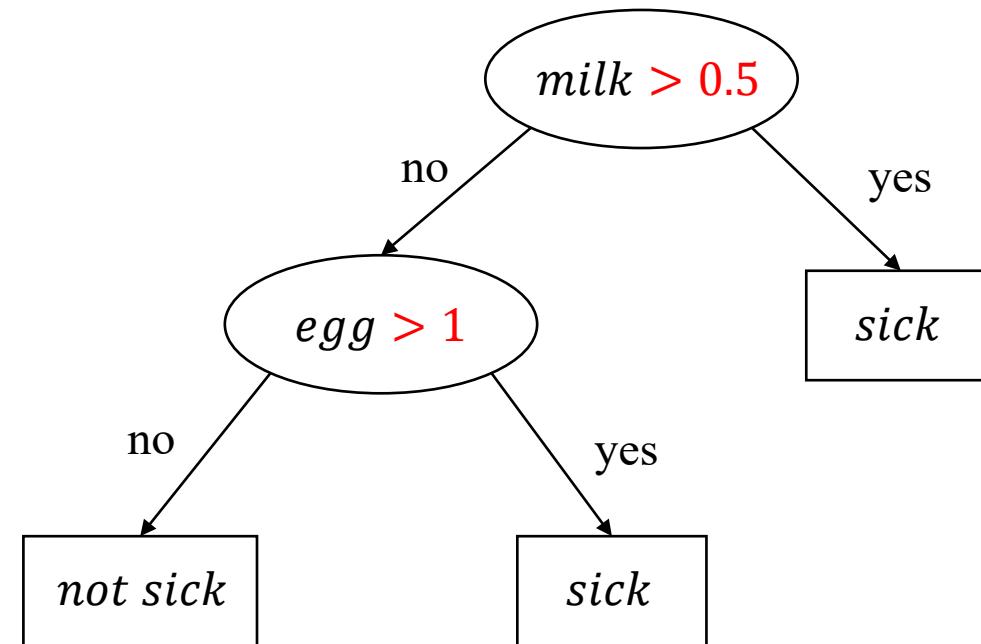
- The previous restaurant example has only **discrete features**, how about **real-value features**?
- If you frequently start getting an upset stomach and suspect an adult-onset food allergy. To solve the mystery, you start a food journal

Egg	Milk	Fish	Wheat	Shellfish	Peanuts	...	Sick?
0.0	0.7	0.0	0.3	0.0	0.00	...	1
0.3	0.7	0.0	0.6	0.0	0.01	...	1
0.0	0.0	0.0	0.8	0.0	0.00	...	0
0.3	0.7	1.2	0.0	0.1	0.01	...	1
0.3	0.0	1.2	0.3	0.1	0.01	...	1

Another Example - Food Allergies

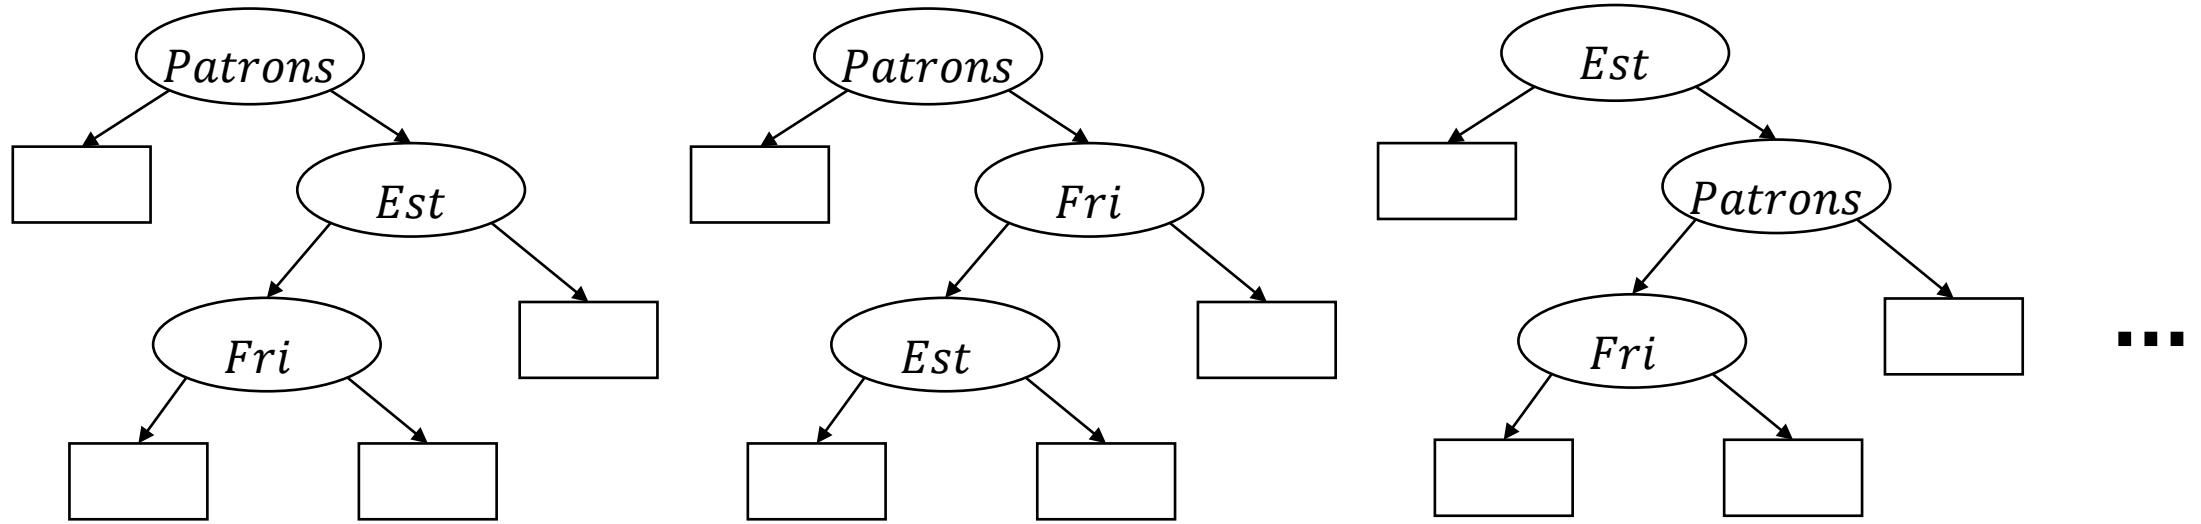
- The splitting rule of the continuous feature contains a threshold.

```
if milk > 0.5 then
    return sick
else
    if egg > 1 then
        return sick
    else
        return not sick
    end
end
```



How do we find the best decision tree?

- The number of possible decision trees is **exponential!!!**



- The number of possible orders to select the attributes is already exponential!
- Learning the smallest decision tree is an **NP-hard problem** (Hyafil & Rivest '76)

Hypothesis Space

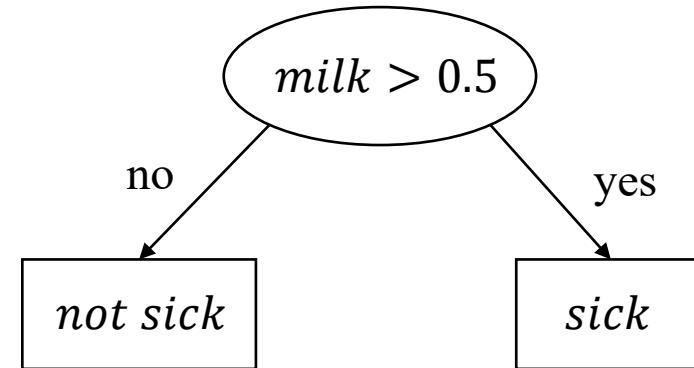
Hypothesis Space - Learning as Search

- Learning can be defined as searching the best hypothesis (e.g., decision tree) for all observed data
 - For decision trees, the hypothesis space are **all possible decision trees** that can be generated for a data set
 - The learner searches through the space and returns the best hypothesis, for decision trees, the tree that potentially best predicts new data
- For a small space, it is possible to test all hypotheses
- When the hypothesis space is large, how do we search the space to find the "best" decision tree?

Decision Tree Learning

Decision Stumps

- The simplest case - "**decision stump**"
 - A decision tree with only **ONE** splitting rule based on thresholding **ONE** feature



- How do we find the best "rule" (feature, threshold, and leaf labels)?
 1. Define a 'score' for the rules
 2. Search for the rule with the best score
- What would you suggest as a score?

Learning a Decision Stump: Accuracy Score

- The most intuitive score: **classification accuracy**
 - How many examples are labelled correctly?
- Computing classification accuracy for ($\text{egg} > 1$):
 - Find most common labels if we use this rule:
 - When ($\text{egg} > 1$), we were "sick" 2 times out of 2
 - When ($\text{egg} \leq 1$), we were "not sick" 3 times out of 4
 - Compute accuracy:
 - The accuracy ("score") of the rule ($\text{egg} > 1$) is 5 times out of 6

Egg	Milk	Fish	...	Sick?
1	0.7	0.0	...	1
2	0.7	0.0	...	1
0	0.0	1.2	...	0
0	0.7	1.2	...	0
2	0.0	1.3	...	1
0	0.0	0.0	...	0

Learning a Decision Stump: Example

- Search for the decision stump maximizing classification score:
 - **Baseline rule** – predict the most common label: this gets 3/6 accuracy
 - If ($milk > 0$) predict "sick" (2/3) else predict "not sick" (2/3): 4/6 accuracy
 - If ($fish > 0$) predict "not sick" (2/3) else predict "sick" (2/3): 4/6 accuracy
 - If ($fish > 1.2$) predict "sick" (1/1) else predict "not sick" (3/5): 4/6 accuracy
 - If ($egg > 0$) predict "sick" (3/3) else predict "not sick" (3/3): 6/6 accuracy
 - If ($egg > 1$) predict "sick" (2/2) else predict "not sick" (3/4): 5/6 accuracy
- Highest-scoring rule: ($egg > 0$), then "sick", else "not sick"
- Questions:
 - Do we need to test the rule ($egg > 3$)?
 - Do we need to test the rule ($egg > 0.5$)?
 - Do we need to test the rule ($egg < 1$)?
- **We only need to test feature thresholds that happen in the data!**

Egg	Milk	Fish	...	Sick?
1	0.7	0.0	...	1
2	0.7	0.0	...	1
0	0.0	1.2	...	0
0	0.7	1.2	...	0
2	0.0	1.3	...	1
0	0.0	0.0	...	0



Decision Tree Learning

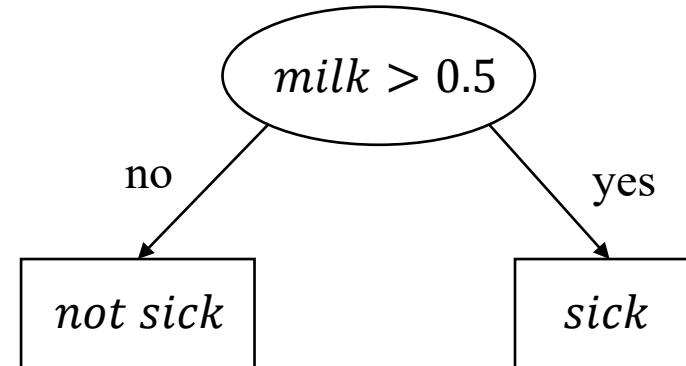
- Decision stumps have only **ONE** rule based on only **ONE** feature
 - Very limited class of models: usually not very accurate for most tasks
- Decision tree learning
 - Recursive stump learning with greedy choice

Example of Greedy Recursive Splitting

Start with the full data set

Egg	Milk	...	Sick?
0	0.7	...	1
1	0.7	...	1
0	0.0	...	0
1	0.6	...	1
1	0.0	...	0
2	0.6	...	1
0	1.0	...	1
2	0.0	...	1
0	0.3	...	0
1	0.6	...	0
2	0.0	...	1

Find the decision stump with the best score



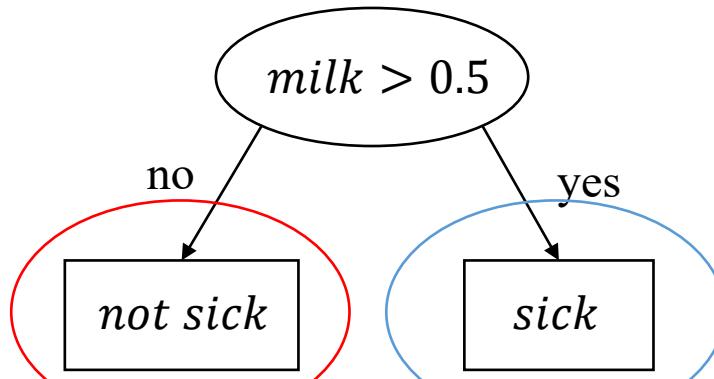
Split into two smaller data sets based on stump

Egg	Milk	...	Sick?
0	0.0	...	0
1	0.0	...	0
2	0.0	...	1
0	0.3	...	0
0	0.0	...	1
2	0.0	...	1

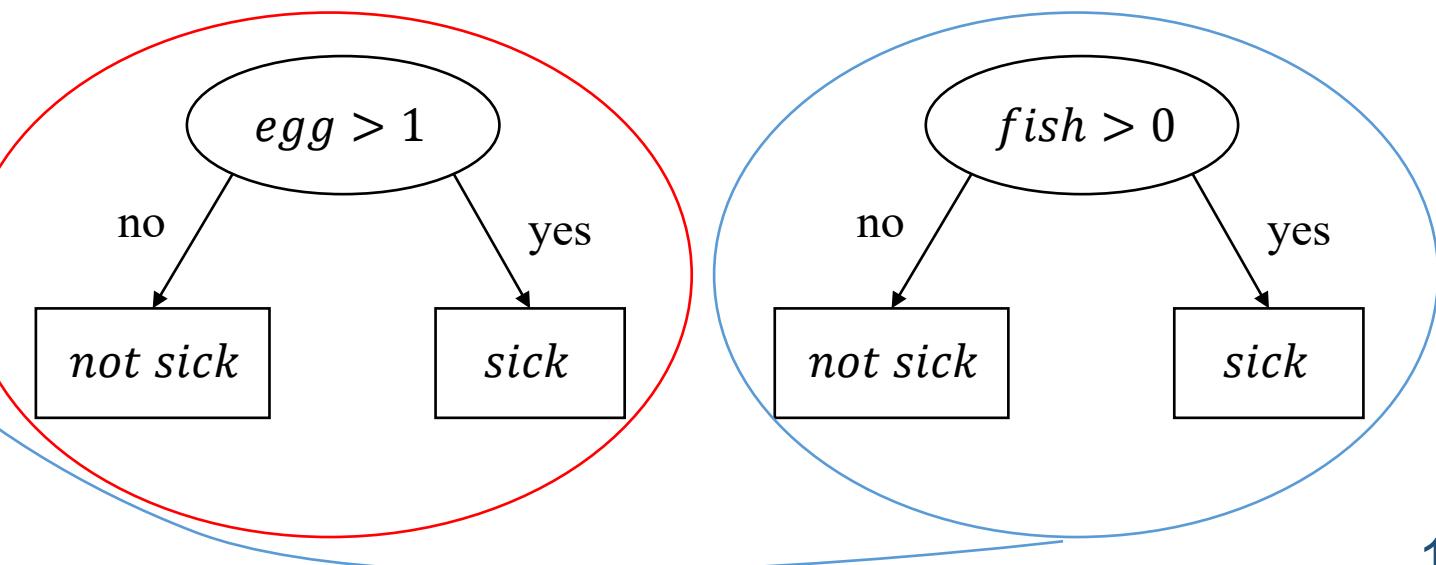
Egg	Milk	...	Sick?
0	0.7	...	1
1	0.7	...	1
1	0.6	...	1
2	0.6	...	1
0	1.0	...	1
1	0.6	...	0
0	1.0	...	1
1	0.6	...	0
1	0.6	...	0

Greedy Recursive Splitting

We now have a decision stump and two data sets



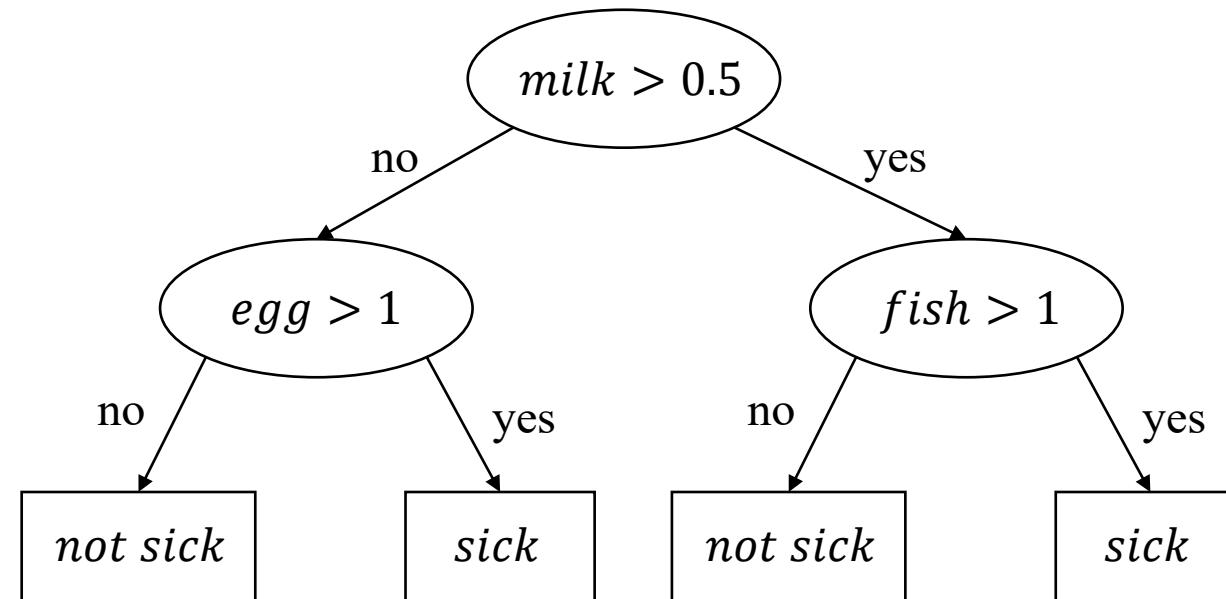
$milk \leq 0.5$					$milk > 0.5$				
Egg	Milk	...	Sick?		Egg	Milk	...	Sick?	
0	0.0	...	0		0	0.7	...	1	
1	0.0	...	0		1	0.7	...	1	
2	0.0	...	1		1	0.6	...	1	
0	0.3	...	0		2	0.6	...	1	
2	0.0	...	1		0	1.0	...	1	



Fit a decision stump to each leaf's data

Then add these stumps to the tree

Greedy Recursive Splitting

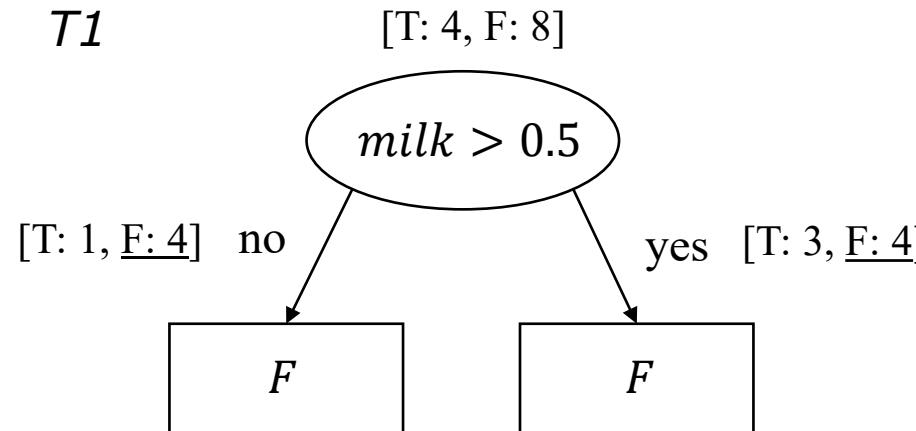


- When to stop splitting?
 1. Can't split a leaf node further, e.g., only 1 example in the leaf node
 2. Leaves have only one label
 3. User-defined maximum depth → Should work, but how to determine the max depth?
 4. Should we stop when accuracy doesn't increase → You might get a shallow tree with low accuracy!
- Multiple criteria can be applied, e.g., 1, 2 and 3

Information Gain

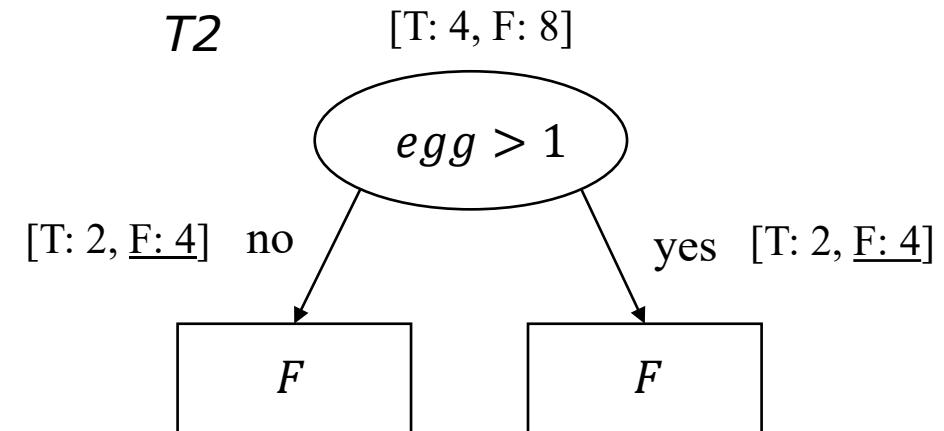
Revisit Accuracy Score...

- Question: Is the accuracy score a good way to choose the right feature to split on?
 - Consider a decision stump (T - Sick, F - Not Sick)



Accuracy before splitting (baseline rule): 2/3

Accuracy of this rule: 2/3



Accuracy before splitting (baseline rule): 2/3

Accuracy of this rule: 2/3

- Both T_1 and T_2 can't improve accuracy!!
 - However, T_1 seems to be a better choice because at least the left branch is more predictable

Choosing a Good Splitting Rule

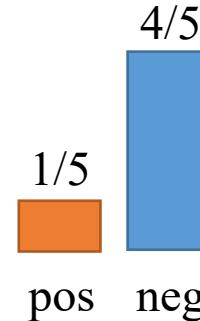
- Intuitively, we want each splitting rule to best distinguish the class labels.
 - The **best case** is to have only one class in each resulting branch (**deterministic**)
 - The **worst case** is all classes are equally probable in each branch (**random**)
- The more deterministic the better!
- How to quantify the uncertainty?

Quantifying Uncertainty

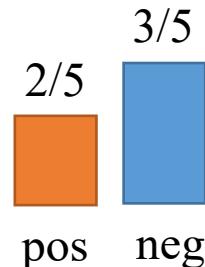
- Entropy: $H(X) = -\sum_{x \in X} p(x) \log_2 p(x)$

- Example: $X = \{pos, neg\}$

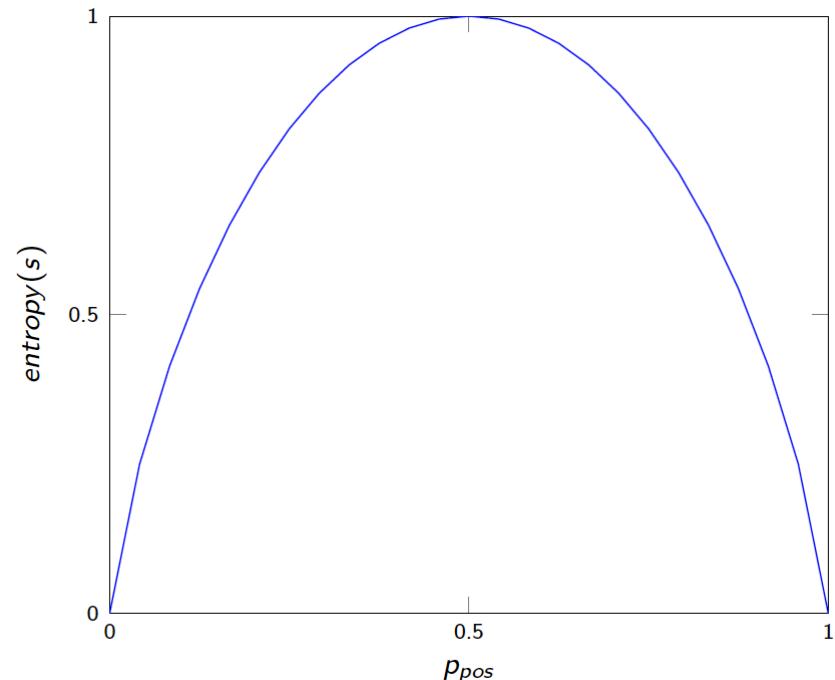
$$H(X) = -p_{pos} \log_2 p_{pos} - p_{neg} \log_2 p_{neg}$$



$$-\frac{1}{5} \log_2 \frac{1}{5} - \frac{4}{5} \log_2 \frac{4}{5} \approx 0.72$$



$$-\frac{2}{5} \log_2 \frac{2}{5} - \frac{3}{5} \log_2 \frac{3}{5} \approx 0.97$$



- High Entropy – more uncertain, e.g., uniform distribution has the highest entropy
- Low Entropy – more certain, e.g., $p_{pos} = 1$ or $p_{neg} = 1$

Information Gain

- Idea: choose the split that decreases the entropy of labels the most
 - The decrease of entropy: “Information Gain” (IG)

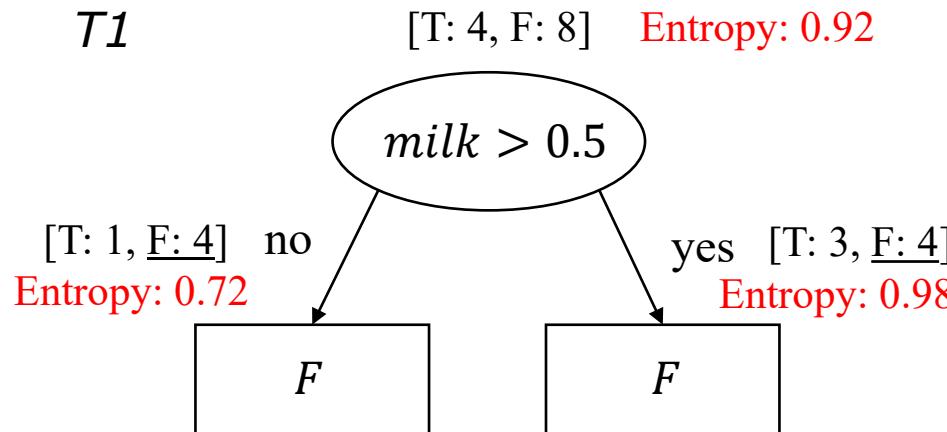
$$IG(S, A) = \underbrace{H(S)}_{\text{entropy before split}} - \sum_{\text{branch} \in B} \overbrace{\frac{|S_{\text{branch}}|}{|S|} \cdot H(S_{\text{branch}})}^{\substack{\text{fraction of examples in each branch} \\ \text{entropy of each branch}}}$$

The expected entropy of all branches

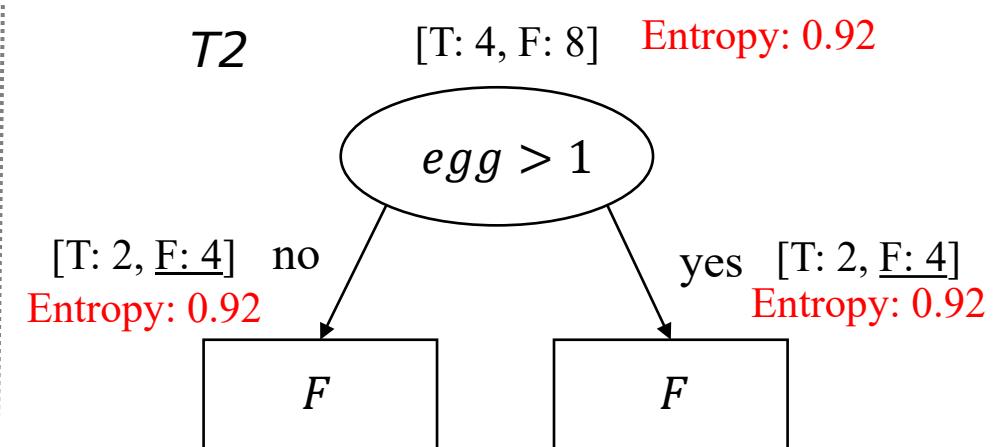
- Information gain for baseline rule (majority class) is 0
- Information gain is large if labels are much "more predictable" ("less random") in next layer

Information Gain: Example

- Consider a decision stump (T - Sick, F - Not Sick)
 - The two examples below have the same accuracy
 - How about entropy?



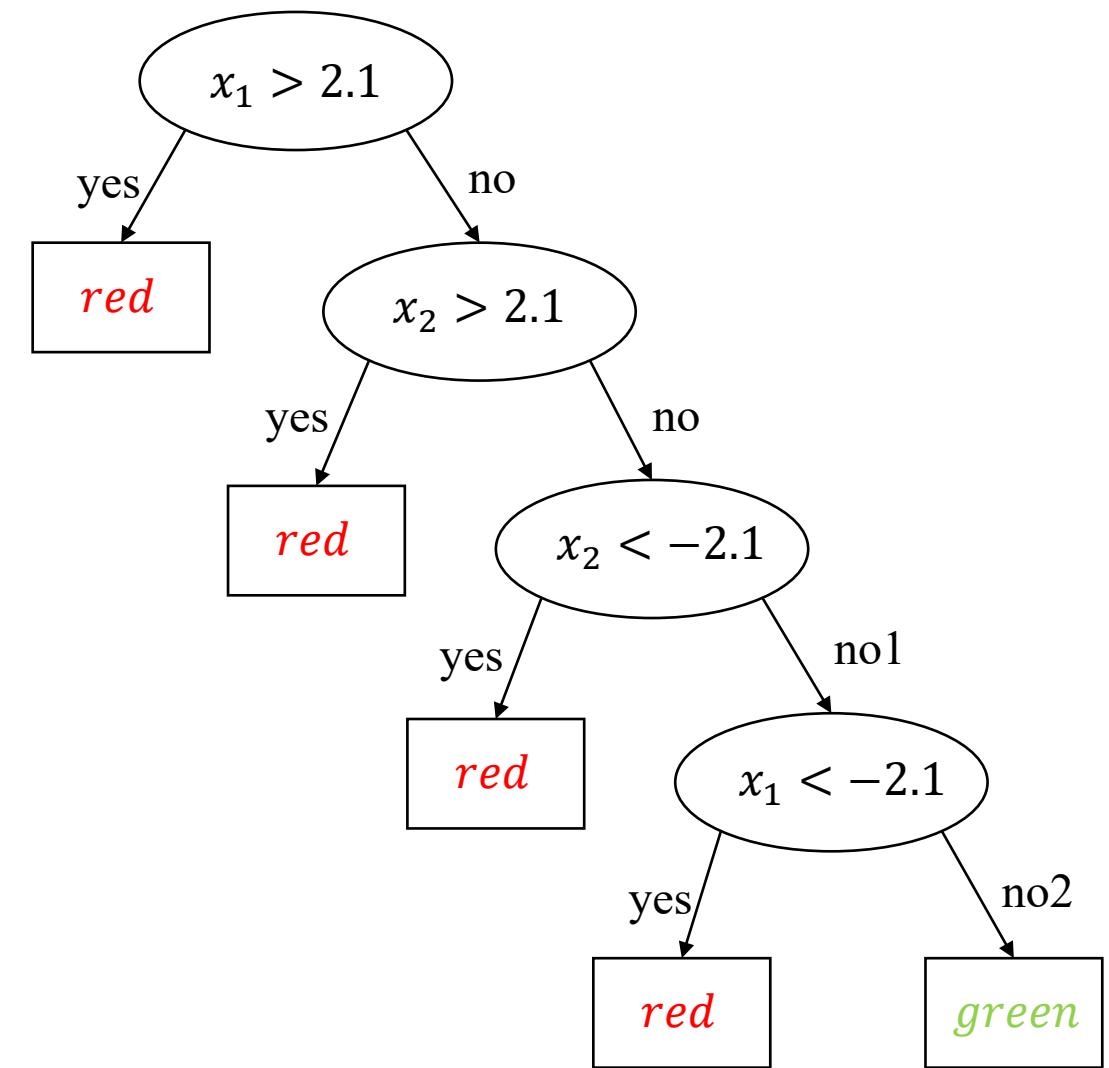
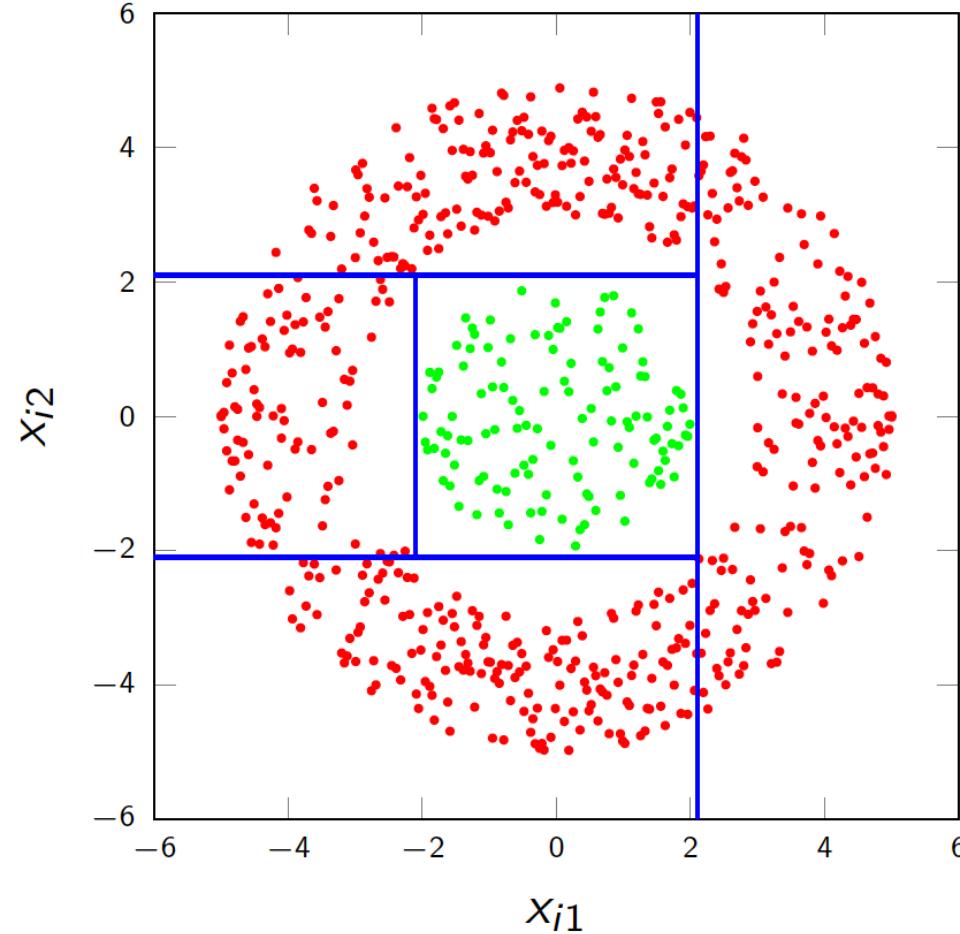
Entropy before splitting (baseline rule): 0.92
 IG of this rule: $0.92 - (5/12)*0.72 - (7/12)*0.98 = 0.048$



Entropy before splitting (baseline rule): 0.92
 IG of this rule: $0.92 - 0.5*0.92 - 0.5*0.92 = 0$

- T1* has a positive information gain because its left branch is more predictable for the negative examples

Decision tree as feature space partitioning



A decision tree partitions the feature space along feature axis!



Decision Trees

- Advantages:
 - Easy to implement
 - Interpretable
 - Learning is fast prediction is very fast
 - Can elegantly handle a small number missing values during training
- Disadvantages
 - Hard to find optimal set of rules
 - Greedy splitting often not accurate, requires very deep trees

Supervised Learning

Supervised Learning

- We can formulate this as a supervised learning problem

Example	Features								Class labels
	Alt	Bar	Fri	Patrons	Price	Res	Type	Est	
x_1	Yes	No	No	Some	\$\$\$	Yes	French	Short	Yes
x_2	Yes	No	No	Full	\$	No	Thai	Mid	No
x_3	No	Yes	No	Some	\$	No	Burger	Short	Yes
x_4	Yes	No	Yes	Full	\$	No	Thai	Mid	Yes
...

⇒

WillWait
Yes
No
Yes
Yes
...

- The input for an **example** (e.g., x_1) is a set of **features** (Alt, Bar, ...)
- The output is a target **class label** (Yes or No)
- Supervised learning:**
 - Use data to find a model that outputs the right label based on the features
 - The model should be able to predict with arbitrary **new feature combinations**.

Supervised Learning Formulation

$$X = \begin{bmatrix} \text{Alt} & \text{Bar} & \text{Fri} & \text{Patrons} & \text{Price} & \text{Res} & \text{Type} & \text{Est} \\ 1 & 0 & 0 & 0 & 3 & 1 & 0 & 0 \\ 1 & 0 & 0 & 1 & 1 & 0 & 1 & 1 \\ 0 & 1 & 0 & 0 & 1 & 0 & 2 & 0 \\ 1 & 0 & 1 & 1 & 1 & 0 & 1 & 1 \\ 1 & 0 & 1 & 1 & 3 & 1 & 0 & 2 \\ 0 & 1 & 0 & 0 & 2 & 1 & 3 & 0 \end{bmatrix} \quad d \quad n \quad y = \begin{bmatrix} 1 \\ 0 \\ 1 \\ 1 \\ 0 \\ 1 \end{bmatrix}$$

- If the attributes are categorical, we can represent the input as a matrix by encoding the input attributes into numeric values:
 - Binary values: Alt: (Yes → 1, No → 0), Bar: (Yes → 1, No → 0)
 - Multiple outcomes: Price: (\$→1, \$\$→2, \$\$\$→3), Type: (French→0, Thai→1, Burger→2, Italian→3)

Supervised Learning Formulation

$$X = \begin{bmatrix} \text{Alt} & \text{Bar} & \text{Fri} & \text{Patrons} & \text{Price} & \text{Res} & \text{Type} & \text{Est} \\ \hline 1 & 0 & 0 & 0 & 3 & 1 & 0 & 0 \\ 1 & 0 & 0 & 1 & 1 & 0 & 1 & 1 \\ 0 & 1 & 0 & 0 & 1 & 0 & 2 & 0 \\ 1 & 0 & 1 & 1 & 1 & 0 & 1 & 1 \\ 1 & 0 & 1 & 1 & 3 & 1 & 0 & 2 \\ 0 & 1 & 0 & 0 & 2 & 1 & 3 & 0 \end{bmatrix} \quad d \quad n$$

WillWait

$$y = \begin{bmatrix} 1 \\ 0 \\ 1 \\ 1 \\ 0 \\ 1 \end{bmatrix}$$

- Feature matrix X has **rows as examples, columns as features**
 - x_{ij} is the **j -th feature** for the **i -th example** (e.g., x_{12} is the Bar attribute of the first example)
 - x_i is the list of all features for the **i -th example**
- Label vector y contains the labels of the examples
 - y_i is the label of the **i -th example** (1 for “wait”, 0 for “do not wait”)

Supervised Learning Notation

$$X = \left[\begin{array}{ccccccc} \text{Egg} & \text{Milk} & \text{Fish} & \text{Wheat} & \text{Shellfish} & \text{Peanuts} & \dots \\ \hline 0.0 & 0.7 & 0.0 & 0.3 & 0.0 & 0.00 & \dots \\ 0.3 & 0.7 & 0.0 & 0.6 & 0.0 & 0.01 & \dots \\ 0.0 & 0.0 & 0.0 & 0.8 & 0.0 & 0.00 & \dots \\ 0.3 & 0.7 & 1.2 & 0.0 & 0.1 & 0.01 & \dots \\ 0.3 & 0.0 & 1.2 & 0.3 & 0.1 & 0.01 & \dots \end{array} \right]_d^n \quad y = \begin{array}{c} \text{Sick?} \\ \hline 1 \\ 1 \\ 0 \\ 1 \\ 1 \end{array}$$

- Training phase
 - Use X and y to find a model (like a decision stump)
- Prediction phase
 - Given an example x_i , use model to predict a label \hat{y}_i ("sick" or "not sick")
- Training error
 - Fraction of times our prediction \hat{y}_i does not equal the true y_i label

Supervised Learning

- General supervised learning problem:
 - Take features of examples and corresponding labels as inputs
 - Find a model that can accurately predict the labels of new examples
- This is the most successful or widely used machine learning technique
 - Spam filtering, optical character recognition, speech recognition, classifying tumours, etc.
- We've talked about categorical labels in the decision tree examples, which is called **classification**. The model is called a **classifier**.
- When the labels are real-value numbers, the problem is called **regression**. The model is called a **regressor**.

Unsupervised Learning

Unsupervised Learning

- Supervised learning:
 - We have features x_i and **class labels y_i**
 - Find a mapping from x_i to y_i
- Unsupervised learning
 - We only have x_i values, but **no explicit labels**
 - Understand the underlying patterns (e.g., clusters)
- Some unsupervised learning tasks
 - Clustering: What types of x_i are there?
 - Association rules: Which x_j occur together?
 - Outlier detection: Is this a 'normal' x_i ?
 - Similarity search: Which examples look like this x_i ?
 - Latent-factors: What 'parts' are the x_i made from?
 - Ranking: Which are the most important x_i ?



Summary

- **Decision trees:** predicting a label using a sequence of simple rules
- **Decision stumps:** simple decision trees with only one splitting rule
- **Greedy recursive splitting:** uses a sequence of decision stumps to grow a tree
 - Very fast and interpretable, but not always the most accurate
- **Information gain:** splitting score based on decreasing entropy
- **Supervised learning v.s. Unsupervised Learning**
 - Supervised learning: finding a mapping from input features to class labels
 - Unsupervised learning: finding patterns within data without explicit labels