

Fundamentals of Learning

Kaiqi Zhao
The University of Auckland

Slides are partially based on the materials from the University of British Columbia

Content

- Training Error and Test Error
- Golden Rule of Machine Learning
- IID
- Fundamental Trade-Off
- Validation Error
- Hyperparameter Tuning
- Classifier Evaluation Metrics

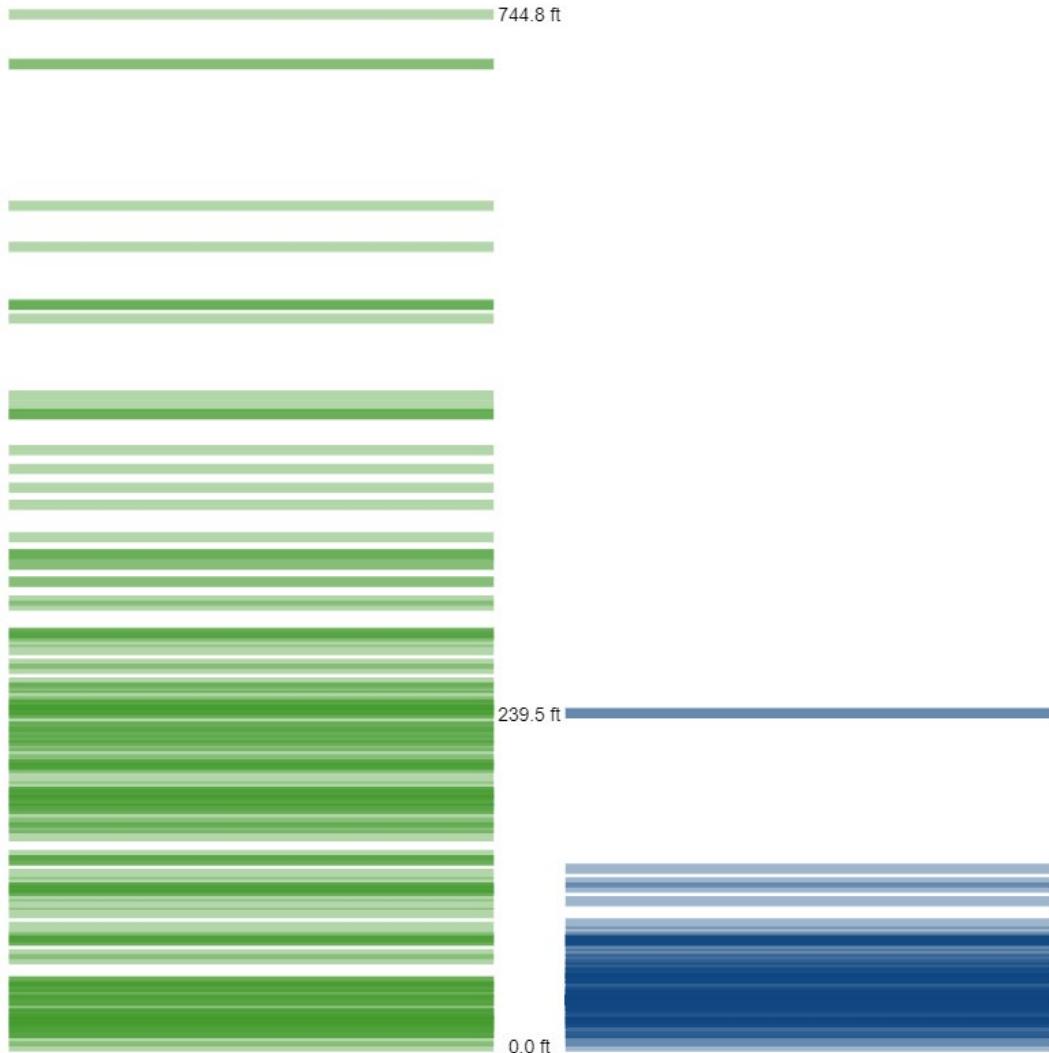
Motivative Example: Determine Home City

- We are given data from 250 homes
- For each home/example, we have these features
 - Elevation
 - Year
 - Bathrooms-Bedrooms
 - Price
 - Square feet
- Goal is to build a program that predicts location (*SF* or *NY*) for homes

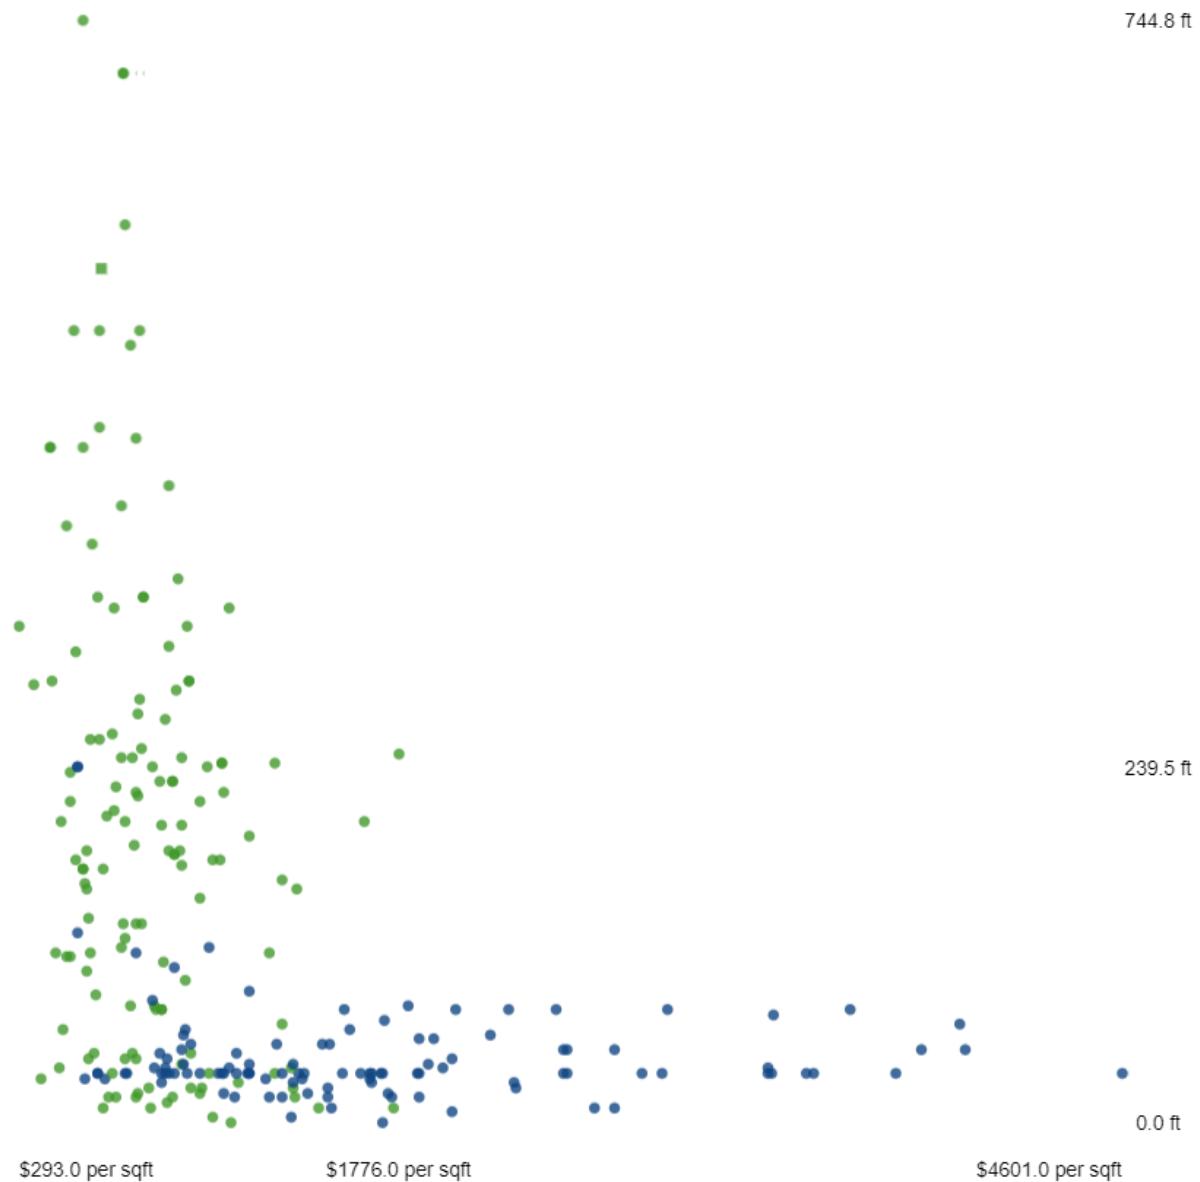
This example and images of it come from: <http://www.r2d3.us/visual-intro-to-machine-learning-part-1>

Elevation

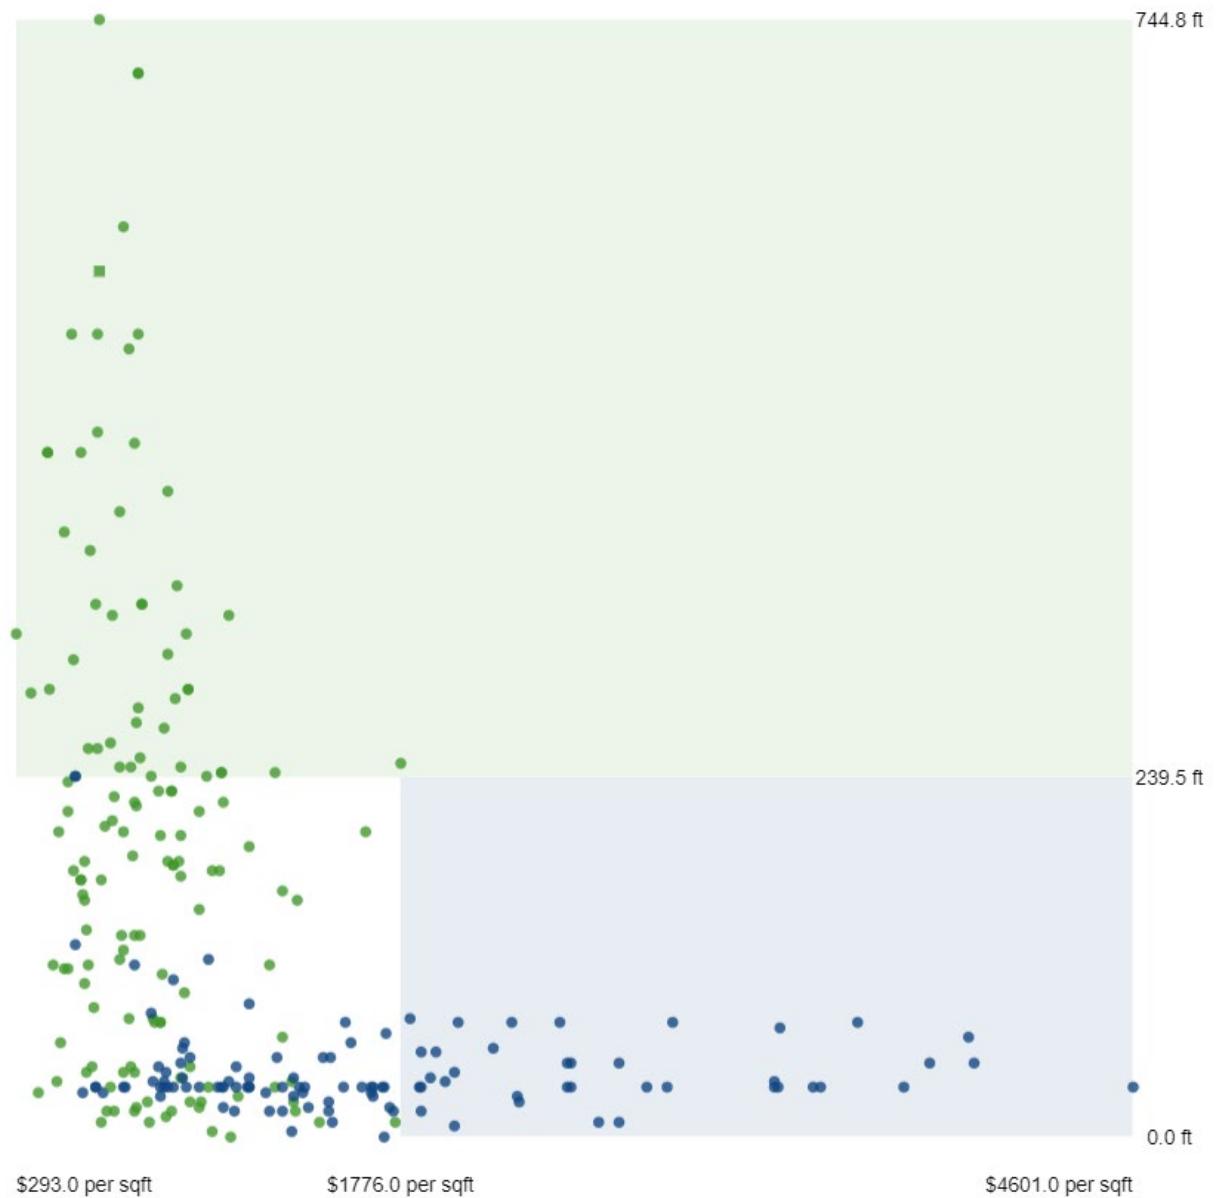
- Elevation (in feet) for San Francisco (SF) and New York (NY).



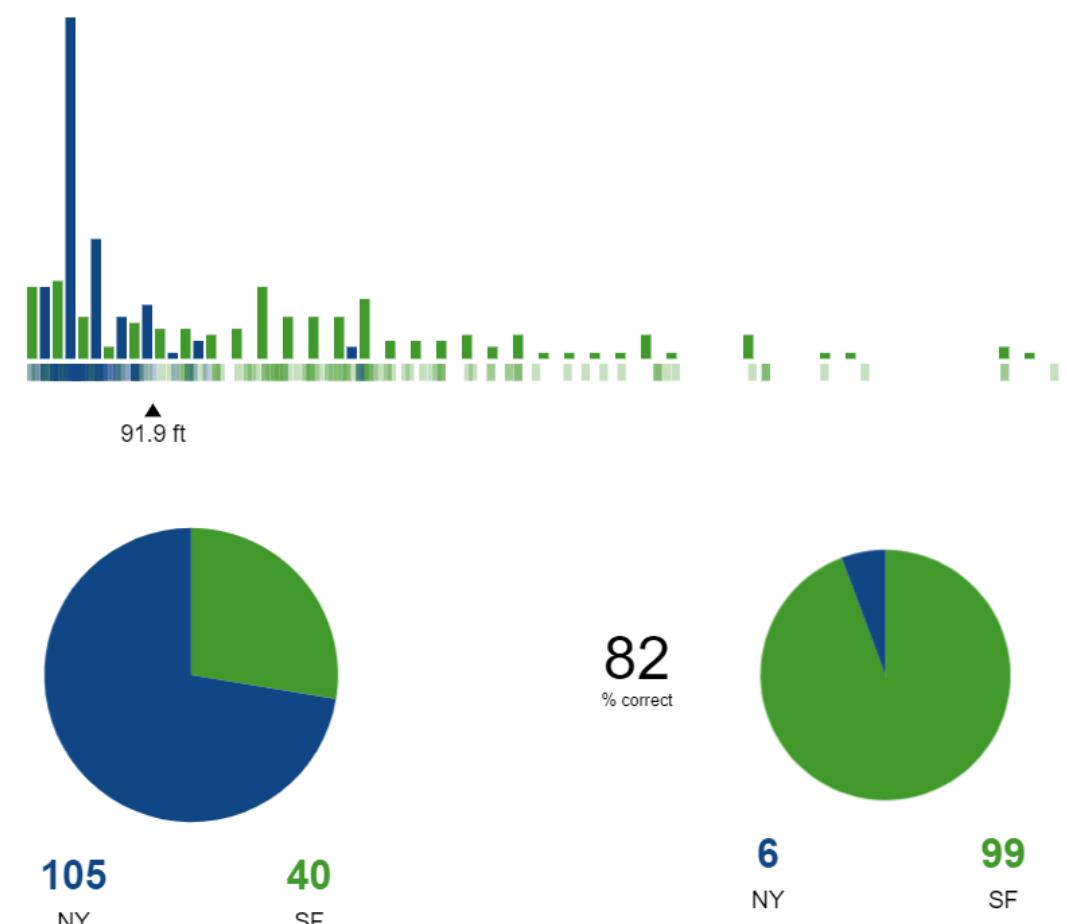
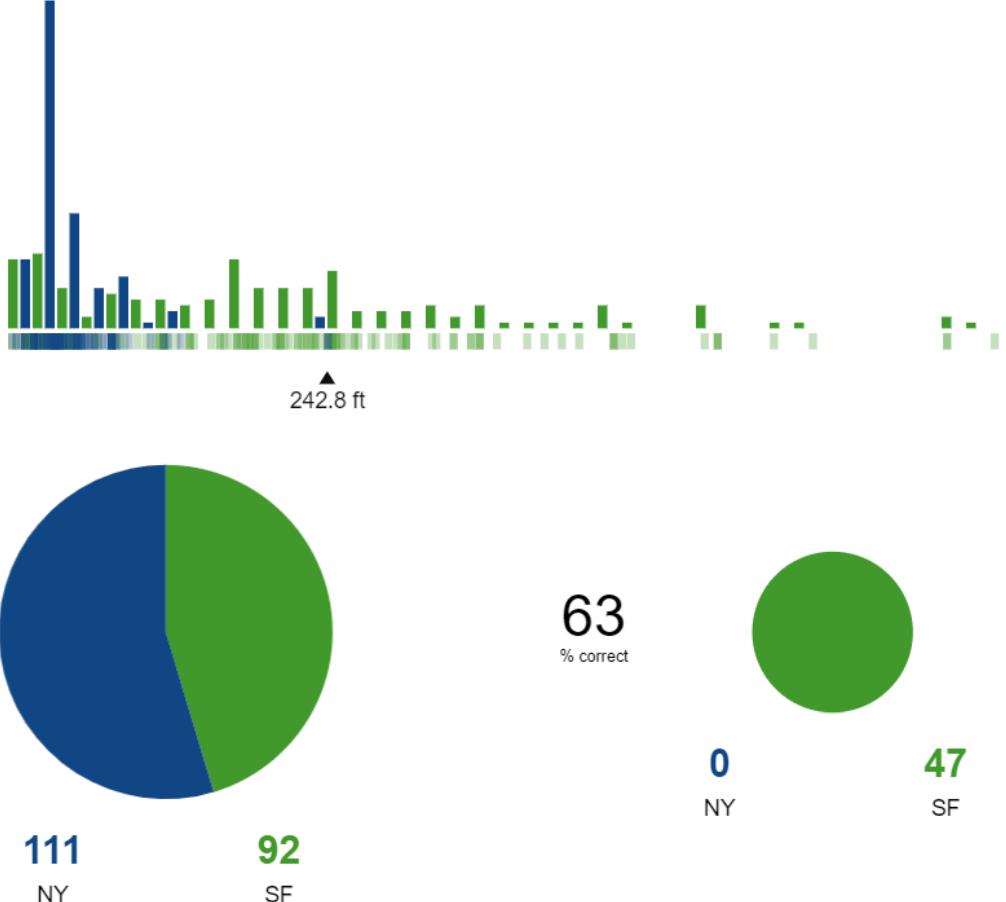
Elevation vs. Price/SqFt



Simple Decision Tree

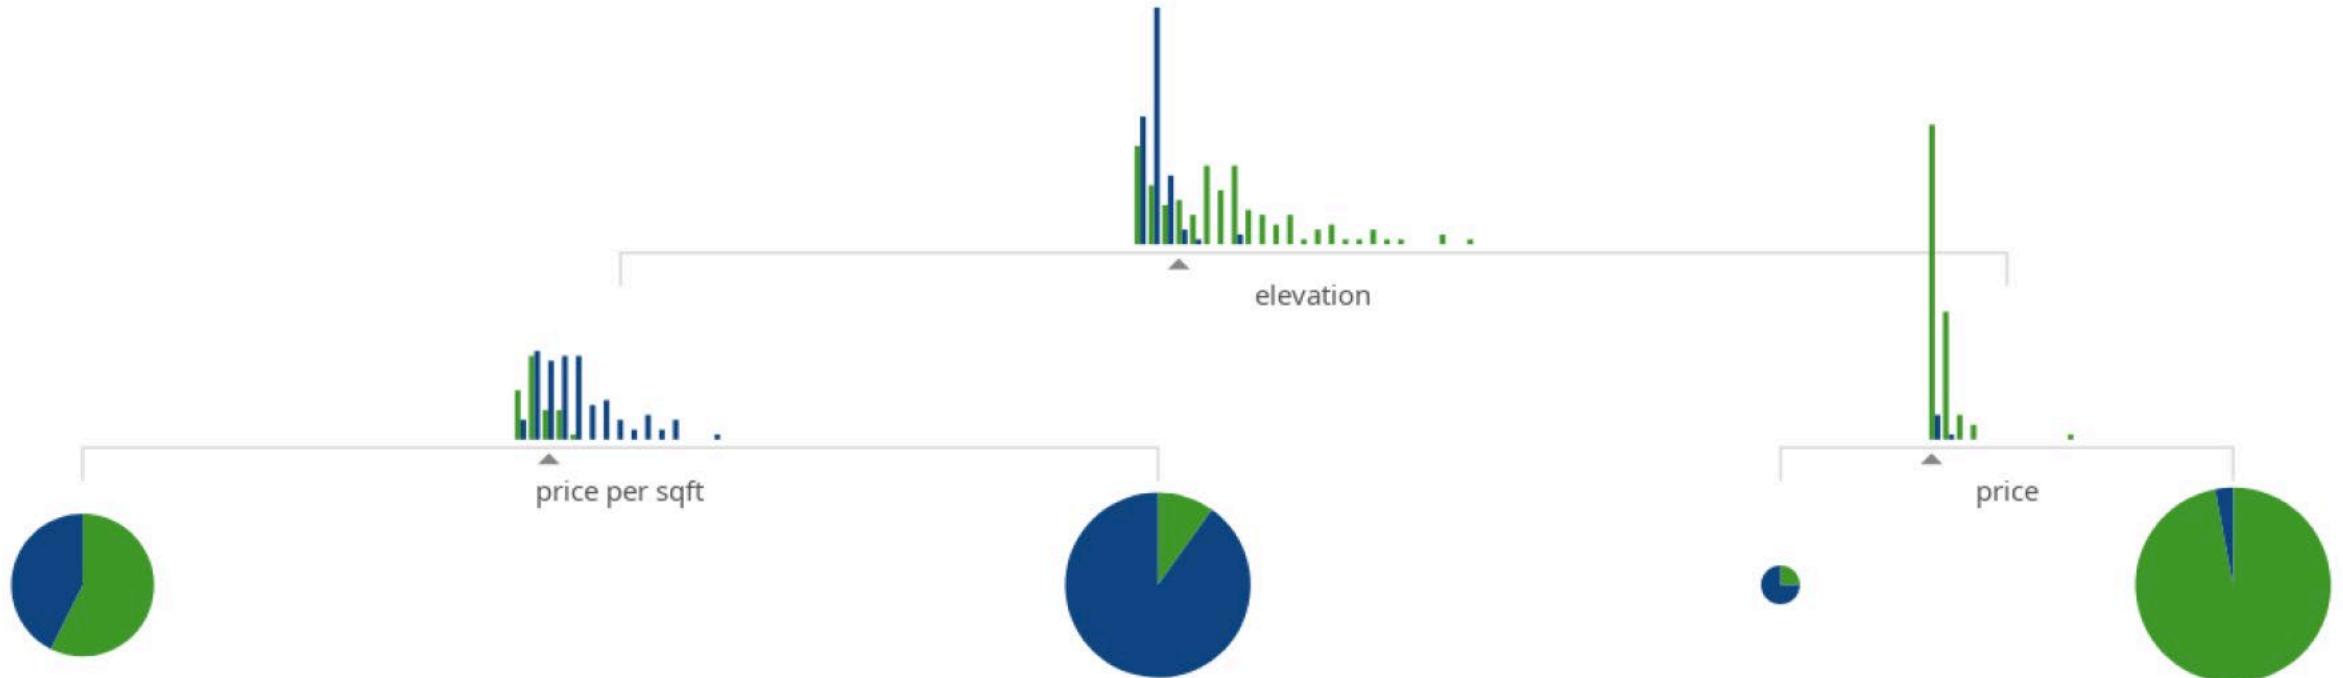


Depth vs. Accuracy

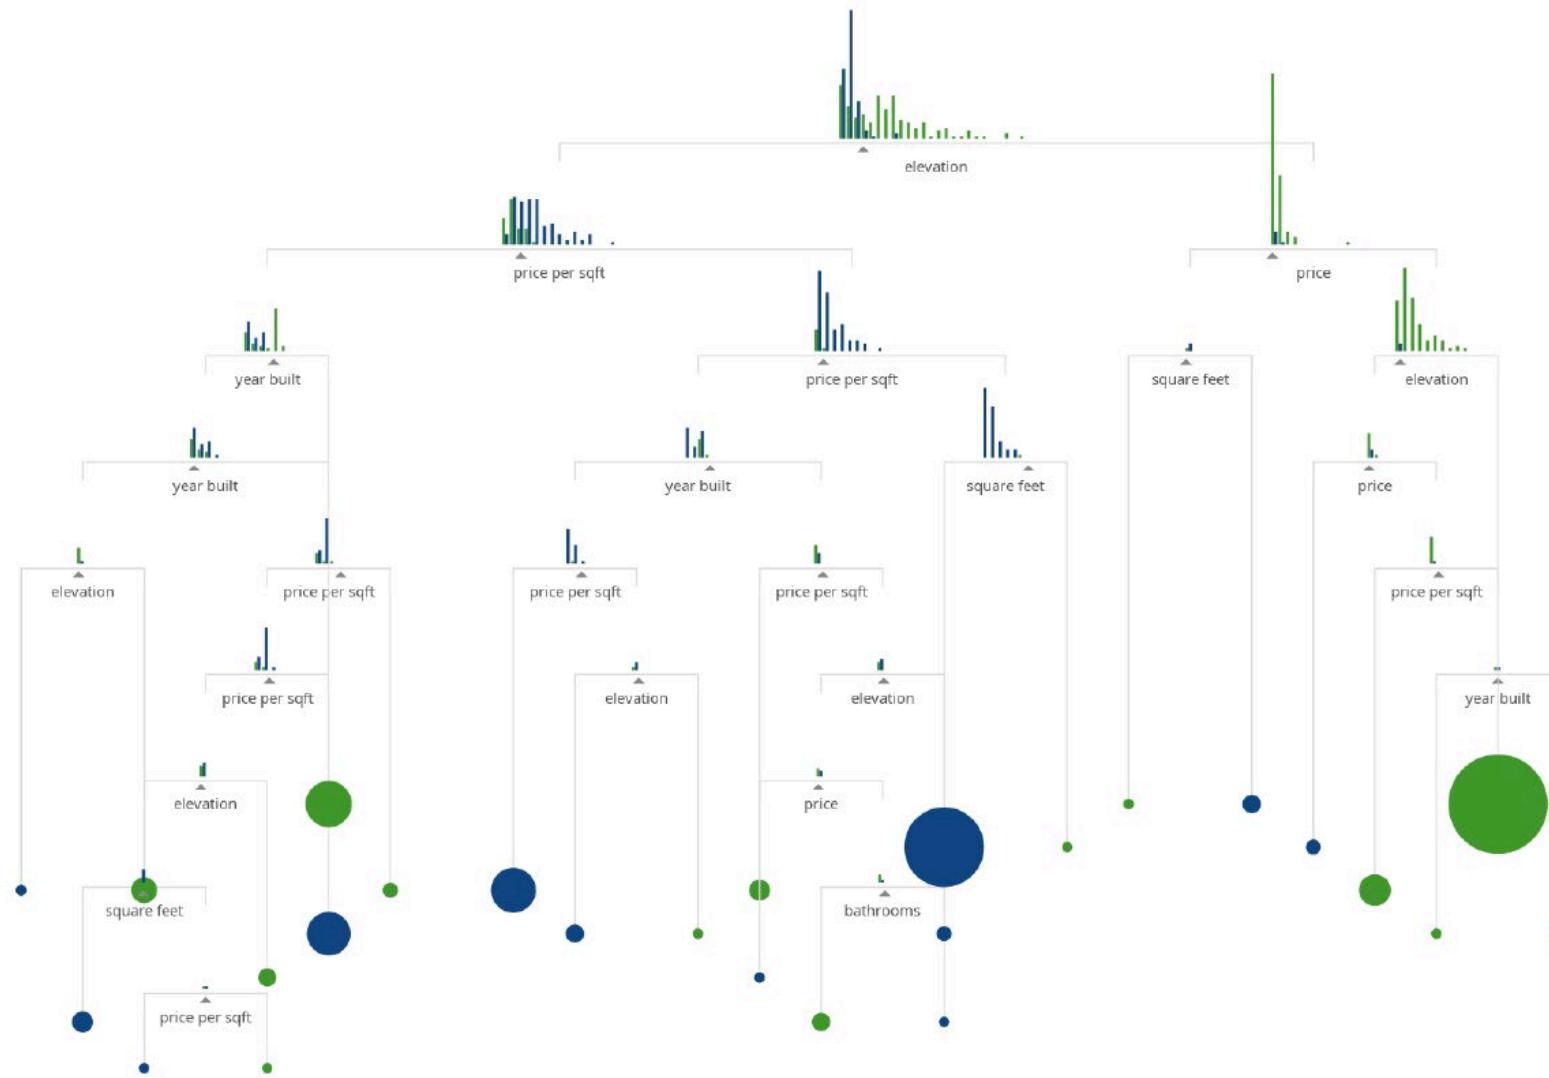


The best split

Depth vs. Accuracy



Depth vs. Accuracy



Training Error and Test Error

Error

- Eventually, we achieved a perfect classification on the data 😊
- With this decision tree, ‘**training accuracy**’ is 1 (training error=0)
 - It perfectly labels the data we used to make the tree



- We are now given features for 242 new homes
- What is the '**testing accuracy**' (or test error) on the new data **NOT used** to make the tree?



- **Overfitting:** lower accuracy on new data
 - Our rules got too specific to the training dataset
 - Some of the "deep" splits only contain a few examples (extreme case - only one example)

Supervised Learning

- We are given training data where we know the labels

	Egg	Milk	Fish	Wheat	Shellfish	Peanuts	...	Sick?
$X =$	0.0	0.7	0.0	0.3	0.0	0.00	...	1
	0.3	0.7	0.0	0.6	0.0	0.01	...	1
	0.0	0.0	0.0	0.8	0.0	0.00	...	0
	0.3	0.7	1.2	0.0	0.1	0.01	...	1
	0.3	0.0	1.2	0.3	0.1	0.01	...	1

$$y =$$

- But there is also testing data we want to label

	Egg	Milk	Fish	Wheat	Shellfish	Peanuts	...	Sick?
$\tilde{X} =$	0.5	0.0	1.0	0.6	2.0	1.0	...	?
	0.0	0.7	0.0	1.0	0.0	0.0	...	$\tilde{y} = ?$
	3.0	1.0	0.0	0.5	0.0	0.0	...	?

Supervised Learning

- Typical supervised learning steps
 1. Build model based on training data X and y (training phase)
 2. Model makes predictions \hat{y} on test data \tilde{X} (testing phase)
- Instead of training error, consider test error:
 - Are predictions \hat{y} similar to true unseen labels \tilde{y} ?

Goal of Machine Learning

- In machine learning:
 - We care more about the test error
- Course analogy:
 - The training error is practice on the exam
 - The test error is the actual exam
 - Goal: do well on actual exam, not only the practice
- Memorization vs learning:
 - Can do well on training data by memorizing it?
 - You have only learned if you can do well in new situations

Golden Rule of Machine Learning

Golden Rule of Machine Learning

- Even though what we care about is test error:
 - The test data **cannot** influence the training phase **in any way**
- We're measuring test error to see how well we do on new data:
 - If used during training, you are not actually measuring the error on **unseen** data, because the test information leaks to your model during training.
 - You can start to overfit if you use it during training
 - Course analogy: You see the exam questions before you attend the exam!

Golden Rule of Machine Learning

- You also shouldn't change the test set to get the result you want
- Note the golden rule applies to hypothesis testing in scientific studies
 - Data that you collect can't influence the hypotheses that you test
- **Extremely common and a major problem**, coming in many forms
 - Collect more data until you coincidentally improve your hypothesis
 - Try different ways to measure the test performance, choose the one that looks best
 - Choose a different type of model / hypothesis after looking at the test data
- If you want to modify your hypotheses, you need to test on new data
- Or at least be aware and honest about this issue when reporting results

Is Learning Possible?

- Does training error say anything about test error?
 - In general, NO: Test data might have nothing to do with training data
 - E.g., "adversary" takes training data and flips all labels.
- In order to learn, we need assumptions:
 - The training and test data need to be related in some way
 - Most common assumption: **independent and identically distributed (IID)**

IID

IID Assumption

- Training/test data is independent and identically distributed (IID) if:
 - All examples come from the same distribution (**identically distributed**)
 - The examples are sampled **independently** (order does not matter)

Age	Job?	City	Rating	Income
23	Yes	Ham	A	22,000.00
23	Yes	Wel	BBB	21,000.00
22	No	Ham	CC	0.00
25	Yes	Akl	AAA	57,000.00

- Examples in terms of cards - which is IID?
 - Pick a card, put it back in the deck, re-shuffle, repeat
 - Pick a card, put it back in the deck, repeat
 - Pick a card, don't put it back, re-shuffle, repeat.

IID in the Food Allergy Example

- Is the food allergy data IID?
 - Do all the examples come from the same distribution?
 - Does the order of the examples matter?
- No!
 - Being sick might depend on what you ate yesterday (not independent)
 - Your eating habits might changed over time (not identically distributed)
- What can we do about this?
 - Just ignore that data isn't IID and hope for the best?
 - For each day, maybe add the features from the previous day?
 - Maybe add time as an extra feature?



Learning Theory

- Why does the IID assumption make learning possible?
 - Patterns in training examples are likely to be the same in test examples
- The IID assumption is rarely true:
 - But it is often a good approximation
 - There are other possible assumptions
- Also, we're assuming IID across examples but not across features
- Learning theory explores how training error is related to test error
- We'll look at a simple example, using this notation:
 - E_{train} is the error on training data
 - E_{test} is the error on testing data

Fundamental Trade-Off

Fundamental Trade-Off

- The **test error** E_{test} could be decomposed to **training error** E_{train} and the error of using E_{train} to approximate E_{test}
- Start with $E_{test} = E_{test}$, then add and subtract E_{train} on the right:

$$E_{test} = \underbrace{(E_{test} - E_{train})}_{\text{test error}} + \underbrace{E_{train}}_{\text{approximation error}} + \underbrace{E_{train}}_{\text{training error}}$$

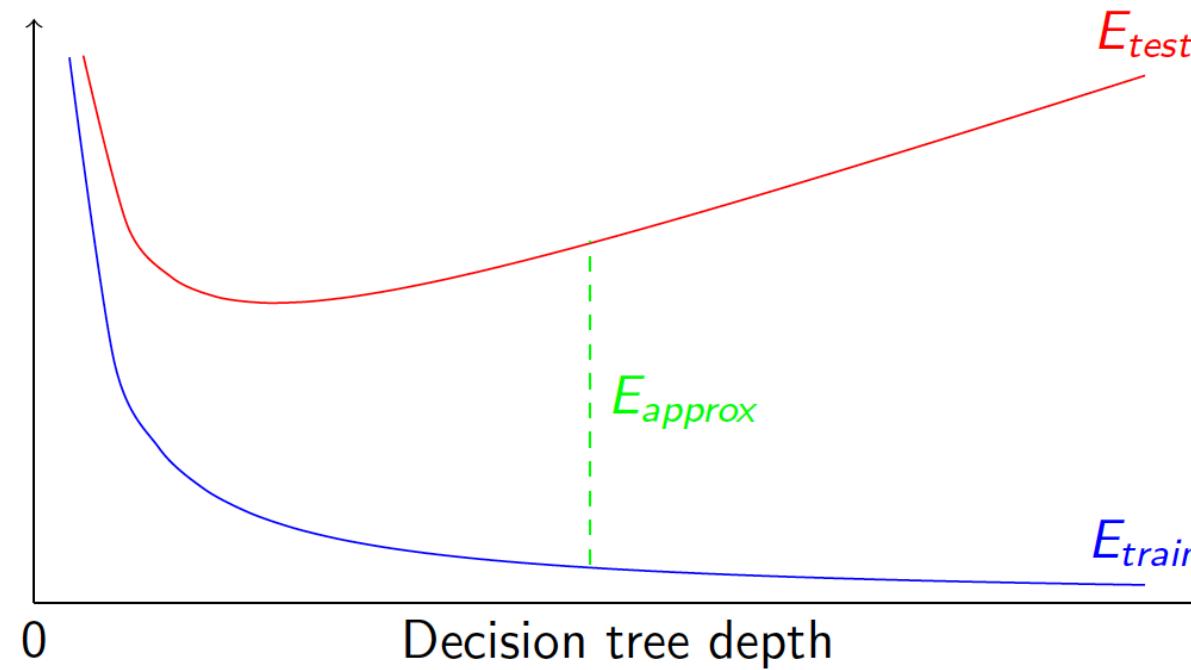
- How does this help?
 - If $E_{approx} = E_{test} - E_{train}$ is small, then E_{train} is a good approximation to E_{test}
- What does E_{approx} ("amount of overfitting") depend on?
 - It tends to get smaller as the training data set gets larger
 - It tends to grow as model get more "complicated"

Fundamental Trade-Off

- This leads to a **fundamental trade-off**:
 1. E_{train} : how small you can make the training error vs.
 2. E_{approx} : how well training error approximates the test error
- Simple models (like decision stumps):
 - E_{approx} is low (not very sensitive to training set)
 - But E_{train} might be high
- Complex models (like deep decision trees):
 - E_{train} can be low
 - But E_{approx} might be high (very sensitive to training set).

Fundamental Trade-Off

- Training error vs. test error for choosing tree depth:
 - Training error is high for low depth (**underfitting**)
 - Training error gets better with depth
 - Test error initially goes down, but eventually increases (**overfitting**)



Refined Fundamental Trade-Off

- Let E_{best} be the irreducible error (lowest possible error for any model)
 - For example, irreducible error for predicting coin flips is 0.5
- Some learning theory results use E_{best} to further decompose E_{test}

$$E_{test} = \underbrace{(E_{test} - E_{train})}_{\text{variance}} + \underbrace{(E_{train} - E_{best})}_{\text{bias}} + \underbrace{E_{best}}_{\text{noise}}$$

- This is similar to the bias-variance decomposition
 - **Variance**: how sensitive we are to training data
 - **Bias**: how low can we make the training error
 - **Noise**: how low can any model make test error



Refined Fundamental Trade-Off

- Decision tree with **high depth**
 - Very likely to fit data well, so **bias** is low
 - But model changes a lot if you change the data, so **variance** is high
- Decision tree with **low depth**
 - Less likely to fit data well, so **bias** is high
 - But model does not change much you change data, so **variance** is low
- And degree does not affect **irreducible error**
 - Irreducible error comes from the best possible model

Bias-Variance Decomposition

- You may have seen bias-variance decomposition
 - Assumes $\tilde{y}_i = \bar{y}_i + \epsilon$, where ϵ has mean 0 and variance σ^2
 - Assumes we have a learner that can take n training examples and use these to make predictions \hat{y}_i
- Expected squared test error in this setting is

$$\underbrace{\mathbb{E}[(\tilde{y}_i - \hat{y}_i)^2]}_{\text{test squared error}} = \underbrace{\mathbb{E}[(\hat{y}_i - \bar{y}_i)^2]}_{\text{bias}} + \underbrace{(\mathbb{E}[\hat{y}_i^2] - \mathbb{E}[\hat{y}_i]^2)}_{\text{variance}} + \underbrace{\sigma^2}_{\text{noise}}$$

- Where expectations are taken over possible training sets of n examples
- Bias is expected error due to having wrong model
- Variance is expected error due to sensitivity to the training set
- Noise (irreducible error) is the best we can hope for given the noise (E_{best})

Bias-Variance vs. Fundamental Trade-Off

- Both decompositions serve the same purpose
 - Trying to evaluate how different factors affect test error
- They both lead to the same 3 conclusions
 1. Simple models can have high E_{train} / bias, low E_{approx} / variance
 2. Complex models can have low E_{train} / bias, high E_{approx} / variance
 3. As you increase n , the number of training examples, E_{approx} / variance goes down

Another Perspective

