

# Bayesian Learning

Kaiqi Zhao  
The University of Auckland

*Slides are partially based on the materials from Mitchel's book and Stanford's NLP lectures*

# Bayes Optimal Classifier

# Bayes Optimal Classifier

- Question: What is the **most probable classification** of the new instance given the training data?
  - Simply applying  $h_{MAP}$  is not the best solution (as one could wrongly think of)
  - Example
    - $H = \{h_1, h_2, h_3\}$ , where  $P(h_1|D) = 0.4$ ,  $P(h_2|D) = P(h_3|D) = 0.3$
    - $h_{MAP} = h_1$
    - Consider a new instance  $x$  encountered, which is classified positive by  $h_1$  and negative by  $h_2, h_3$
    - Taking all hypotheses into account
      - The probability that  $x$  is positive is 0.4 (the posterior probability of  $h_1$ )
      - The probability that  $x$  is negative is 0.6
- ⇒ most probable classification is not the classification result generated by  $h_{MAP}$

# Bayes Optimal Classifier

- The most probable classification is obtained by combining the predictions of all hypotheses, **weighted by their posterior probabilities**

$$P(v_j|D) = \sum_{h_i \in H} P(v_j|h_i)P(h_i|D)$$

where  $P(v_j|D)$  is the probability that the correct classification is  $v_j$

- Bayes optimal classifier

$$\operatorname{argmax}_{v_j \in V} P(v_j|D) = \operatorname{argmax}_{v_j \in V} \sum_{h_i \in H} P(v_j|h_i)P(h_i|D)$$

# Example

$$V = \{\oplus, \ominus\}$$

$$P(h_1|D) = 0.4, P(\ominus|h_1) = 0, P(\oplus|h_1) = 1$$

$$P(h_2|D) = 0.3, P(\ominus|h_2) = 1, P(\oplus|h_2) = 0$$

$$P(h_3|D) = 0.3, P(\ominus|h_3) = 1, P(\oplus|h_3) = 0$$

therefore

$$P(\oplus|D) = \sum_{h_i \in H} P(\oplus|h_i)P(h_i|D) = 0.4$$

$$P(\ominus|D) = \sum_{h_i \in H} P(\ominus|h_i)P(h_i|D) = 0.6$$

and

$$\text{argmax}_{v_j \in \{\oplus, \ominus\}} \sum_{h_i \in H} P(v_j|h_i)P(h_i|D) = \ominus$$

# Naive Bayes Classifier

# Naive Bayes Classifier

- Applies to learning tasks where each instance  $x$  is described by **a conjunction of attribute values  $\langle a_1, a_2, \dots, a_n \rangle$**  and where the target function  $f(x)$  can take any value from some finite set  $V$
- Maximum A Posteriori

$$\begin{aligned}
 v_{MAP} &= \operatorname{argmax}_{v_j \in V} P(v_j | a_1, a_2, \dots, a_n) \\
 &= \operatorname{argmax}_{v_j \in V} \frac{P(a_1, a_2, \dots, a_n | v_j) P(v_j)}{P(a_1, a_2, \dots, a_n)} \\
 &= \operatorname{argmax}_{v_j \in V} P(a_1, a_2, \dots, a_n | v_j) P(v_j)
 \end{aligned}$$

# Naive Bayes Classifier

- $P(v_j)$  can be estimated by counting the frequency of  $v_j$  in  $D$
- $P(a_1, a_2, \dots, a_n | v_j)$  cannot be estimated in this fashion
- Number of these terms is |all possible instances|  $\times$   $|V|$
- Naive Bayes classifier
  - **Assumption: attribute values are conditionally independent given the target value**

$$P(a_1, a_2, \dots, a_n | v_j) = \prod_i P(a_i | v_j)$$

- Hence, number of terms is |distinct attribute values|  $\times$   $|V| + |V|$
- No explicit search through  $H$ , just counting frequencies

Maximum A Posteriori of the prediction target  $\Rightarrow$  Naive Bayes Classifier

$$v_{NB} = \operatorname{argmax}_{v_j \in V} P(v_j) \prod_i P(a_i | v_j)$$

# Naive Bayes for Document Classification

- For a document  $d$  and a class  $c$

$$P(c|d) = \frac{P(d|c)P(c)}{P(d)}$$

- Using MAP:

$$C_{MAP} = \operatorname{argmax}_{c \in C} P(c|d)$$

$$= \operatorname{argmax}_{c \in C} \frac{P(d|c)P(c)}{P(d)}$$

$$= \operatorname{argmax}_{c \in C} P(d|c)P(c)$$

$$= \operatorname{argmax}_{c \in C} P(x_1, x_2, \dots, x_n|c)P(c)$$

# Naive Bayes for Document Classification

$$C_{MAP} = \operatorname{argmax}_{c \in C} P(x_1, x_2, \dots, x_n | c)P(c)$$

- Each document is presented by features  $x_1, x_2, \dots, x_n$
- Let  $X$  be the set of unique feature values (e.g. unique words)
- Assuming the features are independent

$$C_{MAP} = \operatorname{argmax}_{c \in C} P(x_1 | c) \cdot P(x_2 | c) \cdot \dots \cdot P(x_n | c)P(c)$$

$$= \operatorname{argmax}_{c \in C} P(c) \prod_i P(x_i | c)$$

# Naive Bayes for Document Classification

$$C_{MAP} = \operatorname{argmax}_{c \in C} P(c) \prod_i P(x_i | c)$$

- $P(c)$  is easy to estimate – you can count the frequency of the class in your dataset
- How about  $P(x_i | c)$ ?
- Based on the co-occurrence of word  $x_i$  and class  $c$

$$P(x_i | c) = \frac{\operatorname{count}(x_i, c)}{\sum_{x_j \in X} \operatorname{count}(x_j, c)}$$

# Laplace smoothing

$$P(x_i|c) = \frac{\text{count}(x_i, c)}{\sum_{x_j \in X} \text{count}(x_j, c)}$$

- What if  $\sum_{x_j \in X} \text{count}(x_j, c) = 0$ ?
- Solution: Simply add a constant

$$P(x_i|c) = \frac{\text{count}(x_i, c) + 1}{\sum_{x_j \in X} (\text{count}(x_j, c) + 1)}$$

# A very simple Naive Bayes Example

	Doc_ID	Words	class
Training	d1	Kiwi, Sheep, Kiwi	NZ
	d2	Kiwi, Kiwi, Bird	NZ
	d3	Kiwi, Auckland	NZ
	d4	Munich, Oktoberfest, Kiwi	DE
Test	d5	Kiwi, Kiwi, Kiwi, Munich, Oktoberfest	<span style="color:red">NZ</span>

$$P(c) = \frac{N_c}{N}$$

$$P(x|c) = \frac{\text{count}(x, c) + 1}{\sum_{x_j \in X} \text{count}(x_j, c) + |X|}$$

- Priors
 
$$P(NZ) = 3/4$$

$$P(DE) = 1/4$$
- Conditional Probabilities
 
$$P(Kiwi|NZ) = (5 + 1)/(8 + 6) = 3/7$$

$$P(Munich|NZ) = (0 + 1)/(8 + 6) = 1/14$$

$$P(Oktoberfest|NZ) = (0 + 1)/(8 + 6) = 1/14$$

$$P(Kiwi|DE) = (1 + 1)/(3 + 6) = 2/9$$

$$P(Munich|DE) = (1 + 1)/(3 + 6) = 2/9$$

$$P(Oktoberfest|DE) = (1 + 1)/(3 + 6) = 2/9$$
- Choosing a class for  $d_5$ 

$$P(NZ|d_5) \propto 3/4 * (3/7)^3 * 1/14 * 1/14 \approx 0.0003$$

$$P(DE|d_5) \propto 1/4 * (2/9)^3 * 2/9 * 2/9 \approx 0.0001$$

# Bayesian Networks

# Bayesian Networks

- Naive Bayes assumption of conditional independence too restrictive
- But it's intractable without such assumptions
- Bayesian Belief networks describe conditional independence among subsets of variables
- Allows combining prior knowledge about (in)dependencies among variables with observed training data

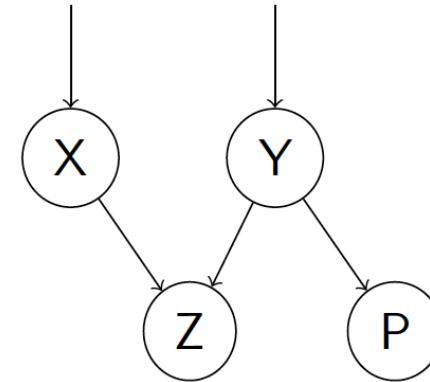
# Bayesian Networks

- A simple, graphical notation for conditional independence assertions and hence for compact specification of full joint distributions
- Syntax:
  - A set of nodes, one per variable
  - A directed, acyclic graph (link  $\approx$  "directly influences")
  - A conditional distribution for each node given its parents:  $P(X_i | Parents(X_i))$
- Conditional distribution represented as a conditional probability table (CPT) giving the distribution of  $X_i$  for each combination of parent values

# Bayesian Networks – A Graphical Model

- A graphical model of relationships
  - Represents dependencies among the variables
    - **Can** be causal dependencies
  - Gives a specification of joint probability distribution:

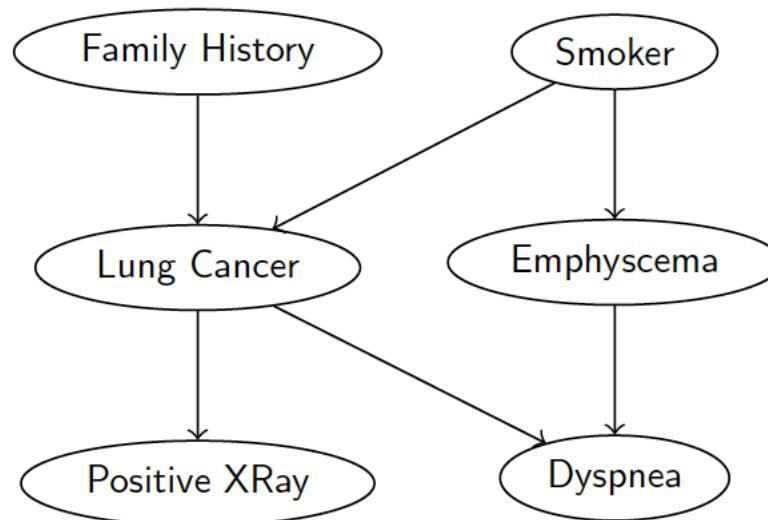
$$P(X_1, \dots, X_n) = \prod_{i=1}^n P(X_i | Parents(X_i))$$



- Nodes: random variables
- Links: dependencies
- $X, Y$  are parents of  $Z$ , and  $Y$  is the parent of  $P$
- Has no loops or cycles

$$P(X, Y, Z, P) = P(Z|X, Y)P(P|Y)P(X)P(Y)$$

# Bayesian Belief Network – Example

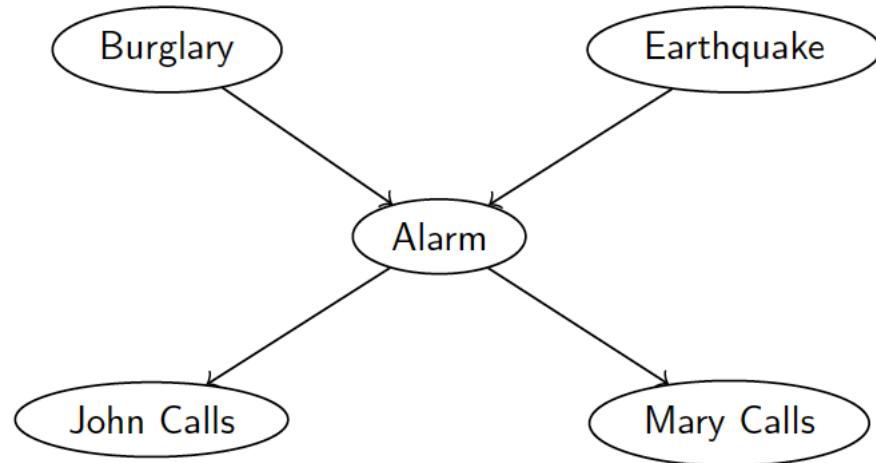


	$FH = T$ $S = T$	$FH = T$ $S = F$	$FH = F$ $S = T$	$FH = F$ $S = F$
$LC = T$	0.8	0.5	0.7	0.1
$LC = F$	0.2	0.5	0.3	0.9

- The conditional probability table for the variable Lung Cancer  $P(LC|FH, S)$
- Shows the conditional probability for each possible combination of its parents

$$P(X_1, \dots, X_n) = \prod_{i=1}^n P(X_i|Parents(X_i))$$

# Bayesian Belief Network – Another Example



$P(B)$	$P(\neg B)$	$P(E)$	$P(\neg E)$		
0.001	0.999	0.002	0.998		
$B$	$E$	$P(A)$	$P(\neg A)$		
T	T	0.95	0.05		
T	F	0.94	0.06		
F	T	0.29	0.71		
F	F	0.001	0.999		
$A$	$P(J)$	$P(\neg J)$	$A$	$P(M)$	$P(\neg M)$
T	0.9	0.1	T	0.7	0.3
F	0.05	0.95	F	0.01	0.99

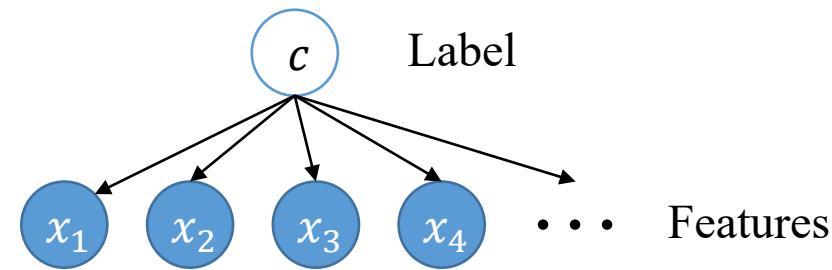
- What is the probability that an alarm has sounded, both Mary and John call, and there was no earthquake and no burglary?

$$P(A, J, M, \neg B, \neg E) = P(J|A) \cdot P(M|A) \cdot P(A|\neg B, \neg E) \cdot P(\neg B) \cdot P(\neg E)$$

$$= 0.9 \cdot 0.7 \cdot 0.001 \cdot 0.999 \cdot 0.998 = 0.00062$$

# Application – Naïve Bayes Classifier

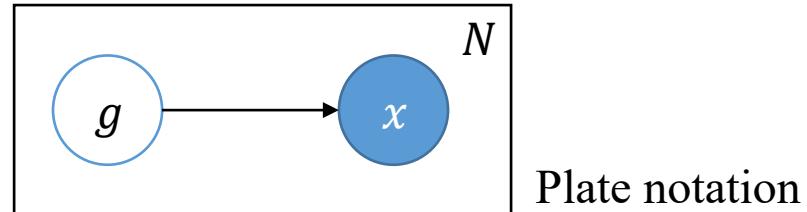
- Naïve Bayes classifier is actually a special case of Bayesian network
  - Variables → features  $x_i$ , class  $c$
  - Graph structure:



$$P(x_1, \dots, x_n, c) = P(c) \prod_i P(x_i | c)$$

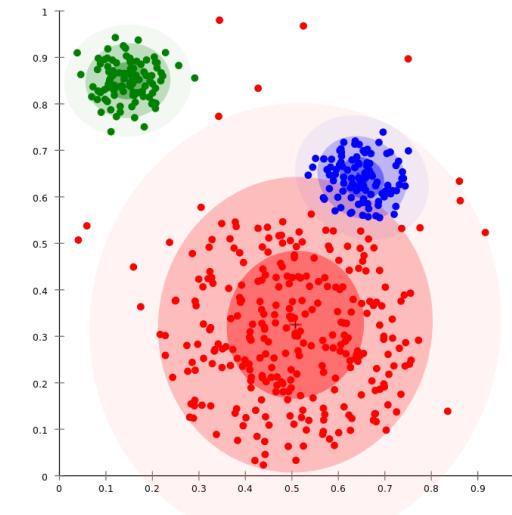
# Application – Gaussian Mixture Model

- Consider a model that clusters data points  $X \in \mathbb{R}^{N \times d}$  in a continuous space into  $K$  groups
  - Variables → data point  $x$ , group  $g$
  - Graph structure:



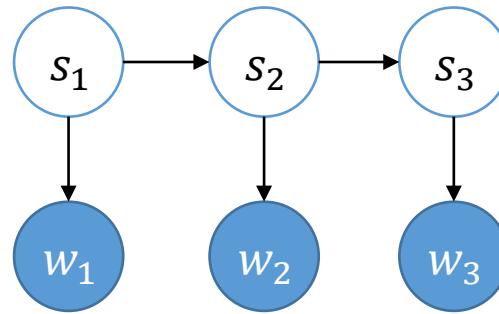
- Probability distribution:
  - $p(g) \sim \text{Categorical}(K)$
  - $p(x|g) \sim \text{Normal}(\mu_g, \Sigma_g)$

$$p(X, G) = \prod_g p(g) \prod_{i=1}^N p(x_i|g_i)$$



# Application – Hidden Markov Model

- Consider a model that predicts tomorrow's weather given the weather condition so far.
  - Variables → state  $s$  (hidden); weather  $w$  (observed)
  - Graph structure:



- Discrete Distribution:
  - $p(s)$
  - $p(s_i|s_{i-1})$
  - $p(w|s)$

$$\begin{aligned}
 p(W, S) &= p(w_1, \dots, w_T, s_1, \dots, s_T) \\
 &= p(s_1) \color{red}{p(w_1|s_1)} p(s_2|s_1) \color{red}{p(w_2|s_2)} \dots p(s_T|s_{T-1}) \color{red}{p(w_T|s_T)} \\
 &= p(s_1) p(s_2|s_1) \dots p(s_T|s_{T-1}) \\
 &\quad \times \color{red}{p(w_1|s_1)} p(w_2|s_2) \dots p(w_T|s_T)
 \end{aligned}$$

# Inference in Bayesian Networks

- We are often interested in the unobservable (hidden) variables, e.g., the hidden states in HMM, the cluster label in GMM.
- If only one variable with unknown value (e.g. target in classification), easy to infer it, e.g., MAP or ML
- In general case, the exact inference is NP-hard, approximations have been introduced (e.g. Monte Carlo-based methods and variational inference)
- In some cases, the structure is unknown, we may need to learn the structure as well!
- Refer to More on Bayesian Networks in Bishop's *Pattern Recognition And Machine Learning*

# Summary

- Bayesian learning relies on Bayes' Theorem
- Bayesian methods can be used to select the most likely hypothesis (MAP/ML) given the data
- Bayesian Learning has multiple roles
  - Provide practical and effective learning algorithms like Naive Bayes
  - Provide a framework
    - For evaluating other learners
    - For analyzing learning
- Bayes optimal classifier combines the predictions of all alternative hypothesis weighted by their posterior probabilities
- Bayesian networks provide a natural representation for conditional independence

# Literature

- Chapter 8 of Bishop's *Pattern Recognition and Machine Learning*
- Chapter 6 of Mitchell's *Machine Learning*