

# Neural Networks




Kaiqi Zhao  
The University of Auckland

*Slides are partially based on the materials from Stanford University and UC Berkeley*

# Regularization

# Overfitting

- Recall the fundamental trade-off



- Neural networks are easy to overfit because of its model complexity and non-linearity

# Regularization

- Regularization is a technique to mitigate the overfitting problem
- In its simplest form, if we have a **minimization** problem:

$$\min_w J(w) = \min_w \sum_{i=1}^m \mathcal{L}(\hat{y}_i, y_i; w)$$

Where  $\mathcal{L}(\hat{y}_i, y_i; w)$  is the loss for each instance. For classification, it could be cross-entropy, for regression, it could be sum of square errors.

- Regularization is to add a penalty  **$R(w)$**  to large values of  $w$

$$\min_w J_R(w) = \min_w \sum_{i=1}^m \mathcal{L}(\hat{y}_i, y_i; w) + \lambda R(w)$$


The hyperparameter  $\lambda$  controls the strengths of the penalty

# Regularization

- Regularization is to add a penalty  $R(w)$  to large values of  $w$


$$\min_w J_R(w) = \min_w \sum_{i=1}^m \mathcal{L}(\hat{y}_i, y_i; w) + \lambda R(w)$$

- L1 regularization:
- $R(w) = \sum_j |w_j|$



L1 is often sparser


- L2 regularization:
- $R(w) = \|w\|^2 = w^T w = \sum_j w_j^2$



[Image from <https://explained.ai/regularization/L1vsL2.html>] 73



# Why Regularization Reduce Overfitting?

- The general idea of regularization is to keep the weights small
- Why does it help?
  - When the weights get close to zero, you can think of some of the neurons are deactivated. This simplifies your model.




# Why Regularization Reduce Overfitting?

- The general idea of regularization is to keep the weights small
- Why does it help?
  - When the weights get close to zero, you can think of some of the neurons are deactivated. This simplifies your model.
  - For some activation function like tanh and sigmoid, small weights around zero makes it closer to a (simple) linear model.




# Other Regularization Techniques

- Dropout – not to rely on a few features, spread out the weights
  - In the **training phase**, randomly deactivate some of the nodes in each instance/batch/mini-batch
  - In the **test phase**, we use the whole network without dropout.



# Other Regularization Techniques

- Dropout
  - Randomly deactivate some of the nodes in each instance/batch/mini-batch
  - How to implement dropout?
    - Let's say you want to keep 80% of your neurons at a layer
    - For each neuron at the layer
      - Generate a random number in the range [0,1] and compare it with 0.8
      - If the random number is smaller than 0.8 keep the neuron, otherwise, multiply the output of the neuron with 0 (i.e., deactivation).
      - Inverted dropout - divide the output of each neuron at this layer by 0.8. **Why??**


$$z_i^{(k)} = g\left(\sum_j W_{i,j}^{(k-1,k)} z_j^{(k-1)}\right)$$


- Remind that in the test phase, we use the whole network
- The values for these neurons are 20% less than usual. Rescale it back to make sure both training and test phases have similar scale.

# Other Regularization Techniques

- Dropout - Example

Keep prob = 2/3



Keep/drop randomly

# Other Regularization Techniques

- Data Augmentation
  - Data transformation - rotation, zoom in/out
  - Add small noise





Image source: <https://nanonets.com/blog/data-augmentation-how-to-use-deep-learning-when-you-have-limited-data-part-2/>

# Other Regularization Techniques

- Early stopping – Stop when the validation error gets worse



# Optimization

# SGD

We often use (mini-batch) SGD for training a neural network. Let us consider a generate optimization problem:


$$L(\mathbf{w}) = \frac{1}{m} \sum_{i=1}^m \mathcal{L}_i(x_i, y_i, \mathbf{w}) + \lambda R(\mathbf{w})$$

$$\nabla_{\mathbf{w}} L(\mathbf{w}) = \frac{1}{m} \sum_{i=1}^m \nabla_{\mathbf{w}} \mathcal{L}_i(x_i, y_i, \mathbf{w}) + \lambda \nabla_{\mathbf{w}} R(\mathbf{w})$$

**Note:** To avoid confusion, we will use  $\mathbf{w}$  (bold font) to denote the weight vector, and  $\mathbf{w}[j]$  to denote the j-th weight in  $\mathbf{w}$ .

# Problems with SGD

- What if loss changes quickly in one direction and slowly in another?
- What does gradient descent do?




SGD often takes noisy updates

# Problems with SGD

What if the loss function has a local minima or saddle point?

Zero gradient, gradient descent gets stuck

Saddle points much more common in high dimension



# SGD + Momentum

- **Motivation:** prevent big changes in **gradients**, try to keep some of the momentum from the past

SGD

$$\mathbf{w}_{t+1} = \mathbf{w}_t - \alpha \nabla f(\mathbf{w}_t)$$


SGD+Momentum

$$\mathbf{v}_{t+1} = \rho \mathbf{v}_t + (1 - \rho) \nabla f(\mathbf{w}_t)$$

$$\mathbf{w}_{t+1} = \mathbf{w}_t - \alpha \mathbf{v}_{t+1}$$

Where  $t$  is denotes the  $t$ -th iteration

- Build up “velocity”  $v$  as a running mean of gradients
- $\rho$  controls how much we keep the previous velocity
- Typically,  $\rho = 0.9$  or  $0.99$



# RMSProp

- **Motivation:** adapting **the learning rate** with running mean of squared gradients

$$S_{t+1, \mathbf{w}[j]} = \beta S_{t, \mathbf{w}[j]} + (1 - \beta) (\nabla_{\mathbf{w}[j]} f(\mathbf{w}_t))^2$$

$$\mathbf{w}_{t+1}[j] = \mathbf{w}_t[j] - \frac{\alpha}{\sqrt{S_{t+1, \mathbf{w}[j]} + \epsilon}} \nabla_{\mathbf{w}[j]} f(\mathbf{w}_t)$$


The gradient along a particular weight  $\mathbf{w}[j]$

Constant to avoid zero denominator, typically  $10^{-8}$

- Note that the learning weight is adjusted differently for different weights  $\mathbf{w}[j]$ 
  - The larger gradient along the weight, the slower update is made
  - The smaller gradient along the weight, the faster update is made

# RMSProp

- Example:
  - Let  $w[2]$  be the direction that has larger gradient, and  $w[1]$  be the direction that has smaller gradient
  - The learning rate for  $w[1]$  will be enlarged, while the learning rate of  $w[2]$  will be decreased



# Adam

- **Motivation:** combine momentum and RMSProp

SGD+Momentum

$$\boldsymbol{v}_{t+1} = \rho \boldsymbol{v}_t + (1 - \rho) \nabla f(\boldsymbol{w}_t)$$

$$\boldsymbol{w}_{t+1} = \boldsymbol{w}_t - \alpha \boldsymbol{v}_{t+1}$$

Adam

$$\boldsymbol{v}_{t+1} = \rho \boldsymbol{v}_t + (1 - \rho) \nabla f(\boldsymbol{w}_t)$$

$$S_{t+1,\boldsymbol{w}[j]} = \beta S_{t,\boldsymbol{w}[j]} + (1 - \beta) (\nabla_{\boldsymbol{w}[j]} f(\boldsymbol{w}_t))^2$$

$$\boldsymbol{w}_{t+1}[j] = \boldsymbol{w}_t[j] - \frac{\alpha}{\sqrt{S_{t+1,\boldsymbol{w}[j]} + \epsilon}} \boldsymbol{v}_t[j]$$


RMSProp

$$S_{t+1,\boldsymbol{w}[j]} = \beta S_{t,\boldsymbol{w}[j]} + (1 - \beta) (\nabla_{\boldsymbol{w}[j]} f(\boldsymbol{w}_t))^2$$

$$\boldsymbol{w}_{t+1}[j] = \boldsymbol{w}_t[j] - \frac{\alpha}{\sqrt{S_{t+1,\boldsymbol{w}[j]} + \epsilon}} \nabla_{\boldsymbol{w}[j]} f(\boldsymbol{w}_t)$$


- People tend to use Adam in most cases
- You can set a higher learning rate  $\alpha$ , the algorithm can adjust it accordingly.
- Typically converge faster

SGD, SGD+Momentum, RMSProp, Adam all have learning rate as a hyperparameter.



Q: Which one of these learning rates is best to use?

# SGD, SGD+Momentum, RMSProp, Adam all have learning rate as a hyperparameter.



⇒ Learning rate decay over time!

**step decay:**

e.g. decay learning rate by half every few epochs.



**exponential decay:**

$$\alpha = \alpha_0 e^{-kt}$$

**1/t decay:**

$$\alpha = \alpha_0 / (1 + kt)$$

SGD, SGD+Momentum, RMSProp, Adam all have learning rate as a hyperparameter.



# Summary

- Neural Networks
  - Stacking non-linear layers of neurons to learn features from different levels
  - Logistic regression is a special case of NN with one single layer.
- Training neural networks
  - Applying the chain rule of derivatives → backpropagation
  - Optimization methods:
    - Gradient ascent/gradient descent
    - Batch/mini-batch/stochastic gradient ascent/descent
    - Momentum, RMSProp, Adam
- Activation functions
- Regularization – L1, L2, Dropout, Data Augmentation



# Literature

- Chapter 5 of Bishop's *Pattern Recognition and Machine Learning*
- Chapter 4 of Mitchell's *Machine Learning*