

# COMPSCI762: Introduction to Machine Learning

## Ensembles

Jörg Simon Wicker  
The University of Auckland



**SCIENCE**  
SCHOOL OF COMPUTER SCIENCE  
MACHINE LEARNING

# Admin

- Any questions on the assignment?
- Plan for the remaining weeks

Week	Topic
Week 8	Instance-based Learning & SVMs
<b>Week 9</b>	<b>Ensembles</b>
Week 10	Clustering
Week 11	Association Rules
Week 12	Outlier Detection

# Motivation

# Motivation



# This week we will cover..

## Motivation

## Ensembles

- Averaging

- Random Forests

- AdaBoost

- XGBoost

*Partially based on Slides from University of British Columbia*

# Ensembles

# Ensembles

- Ensembles are classifiers that have classifiers as input
- With great names:
  - Averaging
  - Boosting
  - Bootstrapping
  - Bagging
  - Cascading
  - Random Forests
  - Stacking
- Ensemble methods often have higher accuracy than the input classifiers
- How would you build an ensemble classifier?


# Ensembles

- Remember the fundamental trade-off
  1.  $E_{train}$ : How small you can make the training error  
vs.
  2.  $E_{approx}$ : How well training error approximates the test error
- Goal of ensemble methods is that the ensemble classifier
  - Does much better on one of these than the individual classifiers
  - Doesn't do too much worse on the other
- This suggests two types of ensemble methods
  1. **Boosting**: Improves training error of classifiers with high  $E_{train}$
  2. **Averaging**: Improves approximation error of classifiers with high  $E_{approx}$

# Averaging


# Averaging

- Input to averaging is the predictions of a set of models, for example
  - Decision trees make one prediction
  - Naïve Bayes makes another prediction
  - KNN makes another prediction
- Simple model averaging:
  - Take the mode of the predictions (or average probabilities if probabilistic)



# Stacking

- A common variation of averaging is stacking
  - Fit another classifier that use the predictions as features



- How does the training set of *decision tree #2* look like?
- Averaging or stacking often performs better than individual models
  - Typically used by Kaggle winners
  - E.g., the Netflix \$1M user-rating competition winner was a stacked classifier

## Why does averaging work?

- Consider 3 binary classifiers, each independently correct with probability 0.8
- With simple averaging, the ensemble is correct if we have *at least 2 right*:

$$P(\text{all 3 right}) = 0.8^3 = 0.512$$

$$P(2 \text{ right, 1 wrong}) = 3 * 0.8^2(1 - 0.8) = 0.384$$

$$P(\text{all 3 right}) = 3 * (1 - 0.8)^2 0.8 = 0.096$$

$$P(\text{all 3 right}) = (1 - 0.8)^3 = 0.512$$

- So the ensemble is right with a probability of 0.896 ( $=0.512+0.384$ )
- Note:
  - For averaging to work, classifiers need to be at least somewhat independent
  - You also want the probability of being right to be  $> 0.5$ , otherwise it will do much worse
  - Probabilities also shouldn't be too different (otherwise, it might be better to take most accurate)

## Why does averaging work?

- Consider a set of classifiers that makes these predictions:
  - Classifier 1: spam
  - Classifier 2: spam
  - Classifier 3: spam
  - Classifier 4: not spam
  - Classifier 5: spam
  - Classifier 6: not spam
  - Classifier 7: spam
  - Classifier 8: spam
  - Classifier 9: spam
  - Classifier 10: spam
- If these independently get 80% accuracy, the mode will be close to 100%
  - In practice errors won't be completely independent due to noise in the labels

## Why can averaging work?

- Why can averaging lead to better results?
- Consider classifiers that overfit (like deep decision trees)
  - If they all overfit in exactly the same way, averaging does nothing – Why?
- But if they make **independent** errors
  - Probability that *average* is wrong can be lower than for each classifier
  - Less attention to specific overfitting of each classifier

## Random Forests

## Random Forests

- Random forests average a set of deep decision trees
  - Tend to be one of the best *out of the box* classifiers
  - Often close to the best performance of any method on the first run
  - And predictions are very fast
- Does averaging work if you use trees with the same parameters?
- Do deep decision trees make independent errors?
  - No: with the same training data you'll get the same decision tree
- Two key ingredients in random forests:
  - Bootstrapping
  - Random trees

# Bootstrap Sampling

- Start with a standard deck of 52 cards
  1. Sample a random card, put it back and re-shuffle
  2. Sample a random card, put it back and re-shuffle
  3. Sample a random card, put it back and re-shuffle
  - ...
  - 52 Sample a random card, put it back and re-shuffle
- Make a new deck of the 52 samples

# Bootstrap Sampling

- The new 52-card deck is called a **bootstrap sample**
- Some cards will be missing, and some cards will be duplicated
  - So calculations on the bootstrap sample will give different results than original data
- However, the bootstrap sample roughly maintains trends:
  - Roughly 25% of the cards will be diamonds
  - Roughly 3/13 of the cards will be face cards
  - There will be roughly four 10 cards
- Common use: compute a statistic based on several bootstrap samples
  - Gives you an idea of how the statistic varies as you vary the data

## Random Forest Ingredient 1: Bootstrap

- Bootstrap sample of a list of  $n$  examples
  - A new set of size  $n$  chosen independently with replacement
 

```

forall  $i \in 1, \dots, n$  do
     $j = \text{rand}(1 : n)$ ; // pick a random number from  $\{1, 2, \dots, n\}$ 
     $X_{\text{bootstrap}}[i, :] = X[j, :]$  // use the random sample
        end

```
  - Gives new dataset of  $n$  examples, with some duplicated and some missing
    - For large  $n$ , approximately 63% of original examples are included (see next slide)
- Bagging: using bootstrap samples for ensemble learning – **Bootstrap Aggregating**
  - Generate several bootstrap samples of the examples  $(x_i, y_i)$
  - Fit a classifier to each bootstrap sample
  - At test time, average the predictions

## 0.632 Bootstrapping

- Probability of an arbitrary  $x_i$  being selected in a bootstrap sample

$p(\text{selected at least once in } n \text{ trials})$

$$= 1 - p(\text{not selected in any of } n \text{ trials})$$

$$= 1 - (p(\text{not selected in one trial}))^n \quad //\text{trials are independent}$$

$$= 1 - (1 - \frac{1}{n})^n \quad //\text{prob} = \frac{n-1}{n} \text{ for choosing any of the } n-1 \text{ other samples}$$

$$\approx 1 - \frac{1}{e} \quad //((1 - \frac{1}{n})^n \rightarrow e^{-1} \text{ as } n \rightarrow \infty)$$


$$\approx 0.632$$

## Random Forest Ingredient 2: Random Trees

- For each split in a random tree model
  - Randomly sample a small number of possible features (typically  $\sqrt{d}$ )
  - Only consider these random features when searching for the optimal rule


### ■ Random tree 1:

*sample(milk, oranges)*



### ■ Random tree 1:

*sample(egg, lactose)*



## Random Forest Ingredient 2: Random Trees

- Splits will tend to use different features in different trees
  - They will still overfit, but hopefully errors will be more independent
- So the average tends to have a much lower test error
- Empirically, random forests are one of the *best* classifiers
- Fernandez-Delgado et al. [2014]
  - Compared 179 classifiers on 121 datasets
  - Random forests are most likely to be the best classifier

# AdaBoost

## AdaBoost: Classic Boosting Algorithm

- A classic boosting algorithm for binary classification is AdaBoost
- AdaBoost assumes we have a *base* binary classifier that
  - Is simple enough that it does not overfit much
  - Can obtain  $> 50\%$  weighted accuracy on any dataset.
- Example: decision stumps or low-depth decision trees
  - Easy to modify stumps/trees to use weighted accuracy as score

# AdaBoost: Classic Boosting Algorithm

- Overview of AdaBoost:
  1. Fit a classifier on the training data
  2. Give a higher weight to examples that the classifier got wrong
  3. Fit a classifier on the weighted training data
  4. Go back to 2
    - Weight gets exponentially larger each time you are wrong.
- Final prediction: weighted vote of individual classifier predictions
  - Trees with higher (weighted) accuracy get higher weight

# AdaBoost: Classic Boosting Algorithm

- Are decision stumps a good base classifier?
  - They tend not to overfit
  - Easy to get  $> 50\%$  weighted accuracy
- Base classifiers that don't work:
  - Deep decision trees (no errors to boost)
  - Decision stumps with infogain (does not guarantee  $> 50\%$  weighted accuracy)
  - Weighted logistic regression (does not guarantee  $> 50\%$  weighted accuracy)

# AdaBoost: Classic Boosting Algorithm

- AdaBoost with shallow decision trees gives fast/accurate classifiers
  - Classically viewed as one of the best *off the shelf* classifiers
  - Procedure originally came from ideas in learning theory
- Many attempts to extend theory beyond binary case
  - Led to *gradient boosting*, which is like *gradient descent with trees*
- Modern boosting methods:
  - Look like AdaBoost, but don't necessarily have it as a special case


# XGBoost

# XGBoost: Modern Boosting Algorithm

- Boosting has seen a recent resurgence, partially due to XGBoost:
  - A boosting implementation that allows huge datasets
  - Has been part of many recent winners of Kaggle competitions
- As base classifier, XGBoost uses *regularized regression trees*

# Regression Trees

- Regression trees used in XGBoost
  - Each split is based on 1 feature
  - Each leaf gives a real-valued prediction
  - Example: *How many hours of video games per day?*



(Sorry for this example!)

- We would predict 2.5 hours for a 14-year-old who owns an Xbox

# Regression Trees


- How can we fit a regression tree?
- Simple approach:
  - Predict: at each leaf, predict mean of the training  $y_i$  assigned to the leaf
    - Weight  $w_L$  at leaf  $L$  is set to  $mean(y_i)$  among  $y_i$  at the leaf node
  - Train: set the  $w - L$  values by minimizing the squared error

$$f(w_1, w_2, \dots) = \sum_{i=1}^n (w_{L_i} - y_i)^2$$

- Same speed as fitting decision trees from earlier in the semester
  - Use mean instead of mode, and use squared error instead of accuracy / infogain
- Use greedy strategy for growing tree, as earlier

## Boosted Regression Trees: Prediction

- Consider an ensemble of regression trees
  - For an example  $i$ , they each make a continuous prediction



- In XGBoost, final prediction is the sum of individual predictions

$$\begin{aligned}
 \hat{y}_i &= \hat{y}_{i_1} + \hat{y}_{i_2} + \hat{y}_{i_3} + \dots + \hat{y}_{i_k} \\
 &= 0.2 + (-0.1) + 0.1 + \dots + (-0.01)
 \end{aligned}$$

## Boosted Regression Trees: Prediction

- Notice we are **not using the mean** as we would with random forests
  - What we do instead?
  - In boosting, each tree is **not individually trying to predict the true  $y_i$  value** (we assume they underfit)
  - Instead, **each new tree tries to fix the prediction made by the old trees**, so that sum is  $y_i$

## Boosted Regression Trees: Training

- Consider the following *gradient tree boosting* procedure:
  - $\text{Tree}[1] = \text{fit}(X, y)$
  - $\hat{y} = \text{Tree}[1].\text{predict}(X)$
  - $\text{Tree}[2] = \text{fit}(X, y - \hat{y})$
  - $\hat{y} = \hat{y} + \text{Tree}[2].\text{predict}(X)$
  - $\text{Tree}[3] = \text{fit}(X, y - \hat{y})$
  - $\hat{y} = \hat{y} + \text{Tree}[3].\text{predict}(X)$
  - $\text{Tree}[4] = \text{fit}(X, y - \hat{y})$
  - $\hat{y} = \hat{y} + \text{Tree}[4].\text{predict}(X)$
  - ...
- Each tree is trying to predict *residuals* ( $\hat{y}_i - y_i$ ) of current prediction
  - “True label is 0.9, old prediction is 0.8, so I can improve  $\hat{y}_i$  by predicting 0.1

# Regularized Regression Trees

- Procedure monotonically decreases the training error
  - As long as not all  $w_L = 0$ , each tree decreases training error
- Can it overfit?
  - It can overfit if trees are too deep or you have too many trees
    - To restrict depth, add L0-regularization (stop splitting if  $w_L = 0$ )

$$f(w_1, w_2, \dots) = \sum_{i=1}^n (w_{L_i} - r_i) + \lambda_0 \|w\|_0$$

- Only split if you decrease squared error by  $\lambda_0$ .
- To further fight overfitting, XGBoost also adds L2-regularization of  $w$

$$f(w_1, w_2, \dots) = \sum_{i=1}^n (w_{L_i} - r_i) + \lambda_0 \|w\|_0 + \lambda_2 \|w\|^2$$

## XGBoost Discussion

- Instead of pruning trees if score does not improve, grows full trees
  - And then prunes parts that don't improve score with L0-regularizer added
- Cost of fitting trees in XGBoost is same as usual decision tree cost
  - XGBoost includes a lot of tricks to make this efficient
  - But cannot be done in parallel like random forest – Why?
- In XGBoost, it's the residuals that act like the *weights* in AdaBoost
  - Focuses on decreasing error in examples with large residuals
- How do you maintain efficiency if not using squared error?
  - For non-quadratic losses like logistic, there is no closed-form solution
  - Approximates non-quadratic losses with second-order Taylor expansion
    - Maintains least squares efficiency for other losses (by approximating with quadratic)

# Summary

- Ensembles combine predictions of base classifiers
- Boosting: ensemble methods that improve training error
- XGBoost: modern boosting method based on regression trees
  - Each tree modifies the prediction made by the previous trees

## Discussion

- In which cases would you **not** use Ensembles?
- How would you handle preprocessing in ensembles?
  - How about feature selection / importance?

## Literature

- Chapter 14 in Bishop
- Chapter 11 in Peter Flach's *Machine Learning* (short)



**SCIENCE**  
SCHOOL OF COMPUTER SCIENCE  
MACHINE LEARNING

Thank you for your attention!

<https://ml.auckland.ac.nz>